



Computational Storage Architecture and Programming Model

Version 1.0

Abstract: This SNIA document defines recommended behavior for hardware and software that supports Computational Storage.

This document has been released and approved by the SNIA. The SNIA believes that the ideas, methodologies and technologies described in this document accurately represent the SNIA goals and are appropriate for widespread distribution. Suggestions for revisions should be directed to <https://www.snia.org/feedback/>.

SNIA Standard

August 30, 2022

USAGE

Copyright © 2022 SNIA. All rights reserved. All other trademarks or registered trademarks are the property of their respective owners.

The SNIA hereby grants permission for individuals to use this document for personal use only, and for corporations and other business entities to use this document for internal use only (including internal copying, distribution, and display) provided that:

1. Any text, diagram, chart, table or definition reproduced shall be reproduced in its entirety with no alteration, and,
2. Any document, printed or electronic, in which material from this document (or any portion hereof) is reproduced, shall acknowledge the SNIA copyright on that material, and shall credit the SNIA for granting permission for its reuse.

Other than as explicitly provided above, you may not make any commercial use of this document or any portion thereof, or distribute this document to third parties. All rights not explicitly granted are expressly reserved to SNIA.

Permission to use this document for purposes other than those enumerated above may be requested by e-mailing tcmd@snia.org. Please include the identity of the requesting individual and/or company and a brief description of the purpose, nature, and scope of the requested use.

All code fragments, scripts, data tables, and sample code in this SNIA document are made available under the following license:

BSD 3-Clause Software License

Copyright (c) 2022, The Storage Networking Industry Association.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

* Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.

* Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

* Neither the name of The Storage Networking Industry Association (SNIA) nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

DISCLAIMER

The information contained in this publication is subject to change without notice. The SNIA makes no warranty of any kind with regard to this specification, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose. The SNIA shall not be liable for errors contained herein or for incidental or consequential damages in connection with the furnishing, performance, or use of this specification.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Revision History

Major	Revision	Editor	Date	Comments
0.1	R1	Stephen Bates	02-Feb-2019	Initial conversion of the NVM Programming Model document.
0.1	R2	Stephen Bates	11-Mar-2019	Added initial draft of Scope section with some initial feedback
0.1	R3	Stephen Bates David Slik	02-April-2019	Added diagram to Scope section and introduced “Computational Function” as a term.
0.1	R4	Nick Adams	04-Apr-2019	Updates per Morning Session on 03-Apr-2019
0.1	R5	Scott Shadley	24-Apr-2019	Updates added per F2F meeting in Boston
0.1	R6	Stephen Bates and Bill Martin	21-Aug-2019	Removed resolved comments, added Theory of Operation, added illustrative diagrams, added placeholder for an illustrative example. Also performed a large amount of editorial clean-up. Added forward and removed taxonomy and abstract (approved by TWG on August 28 th 2019).
0.1	R7	Stephen Bates and Bill Martin	XX Oct 2019	Update foreword section. Add references, abbreviations and definitions.
0.1	R8	David Slik	12 Nov 2019	Moved fields into tables, updated diagrams.
0.1	R9	Stephen Bates	14 Nov 2019	Minor editorial changes
0.1	R10	Bill Martin	20 Nov 2019	Incorporated approved dictionary terms, abbreviations, and reference modifications
0.1	R11	Bill Martin	11 Dec 2019	Incorporated NetInt proposal for Video Compression CSS Incorporated ScaleFlux proposal for Database Filter CSS Removed previously accepted changes
0.1	R12	Stephen Bates	15 Dec 2019	Incorporated DellEMC proposal for Hashing and CRC Removed list of CSSes at start of Section 5. Remove reference to how to add new CSSes to Section 5. Update the list of contributors as per the list on the Wiki Resolved and removed comments. Added illustrative examples from Eideticom and Arm/NGD Systems in their current form and started merging them correctly into this document.. Incorporated Seagate proposal for large dataset PCSS
0.1	R13	Stephen Bates	18 Dec 2019	Removed Annex B (Illustrative Examples)
0.3	R1	Stephen Bates	18 Dec 2019	Update author list Slight fix to Philip's PCSS example Add Raymond's FCSS example Minor corrections to document.
0.4	R1	Bill Martin	8 April 2020	Added PCSS on a Large Multi-Device Dataset Added PCSS – Linux Container Example
0.4	R2	Bill Martin	20 May 2020	Added FCSS – Data Deduplication Added PCSS – OPEN CL example
0.4	R3	Bill Martin	29 July 2020	Updated example in B.4 Data Deduplication Example Added Data Compression Illustrative Example Added Data Filtration Illustrative Example
0.4	R4	Bill Martin	5 August 2020	Added Scatter Gather illustrative example
0.5	R1	Bill Martin	10 Aug 2020	Minor editorial cleanup Removed editor's notes in preparation for public review release
0.5	R2	Bill Martin	12 Aug 2020	Modifications created during TWG call on August 12. There are some partial sentences as the discussion was ended due to the end of the call.
0.5	R3	Bill Martin	8 September	Includes editor's comments removed from public review version Includes comments received to date (Intel (Kim) and HPe (Chris)) Changes for the terminology of CSS, FCSS, PCSS, and CSP

0.5	R4	Bill Martin	30 September	Renaming of components (e.g., addition of Computational Storage Engine, Computational Storage Service becomes Computational Storage Function)
0.5	R5	Bill Martin	9 November	Changes from meetings Modified figures as requested Added memory definitions Removed resolved comments
0.5	R6	Bill Martin	10 Dec 2020	Removed change tracking R4 to R5 Fixed figure 1.1 to include CFM Added brief descriptions of CSFM, FDM, AFDM in 4.1 Moved example discover request and response to an annex
0.5	R7	Scott Shadley	16 Dec 2020	Comments from TWG meeting
0.5	R8	Scott Shadley	1 Jan 2021	Comments from TWG meeting
0.5	R9	Bill Martin	28 Jan 2021	Updated direct/indirect model Moved discovery/configuration request/response to annex Other editorial changes from Jan 27, 2021 meeting
0.5	R10	Bill Martin	28 Jan 2021	Removed redline and completed comments
0.5	R11	Bill Martin	10 Feb 2021	Resolved a number of comments from TWG meetings See individual comments with resolutions
0.5	R12	Bill Martin	11 Feb 2021	Removed redline and completed comments
0.5	R13	Bill Martin	16 Feb 2021	Modifications from TWG meeting
0.5	R14	Bill Martin	23 Mar 2021	Updated Large Data Set CSFs Updating Appendix B, sect 5 (Compression) and 6 (data filtering) illustrative examples to align with the CSx/CSE/CSF terminology. (change) Removed all resolved comments
0.5	R15	Bill Martin	23 April 2021	Updates from work group meetings with change bars
0.5	R16	Bill Martin	27 April 2021	Changes to 4.2 and 4.3 as agreed to in TWG meeting Incorporated architecture model changes agreed in TWG meeting Incorporated changes to section 5 as agreed in TWG meeting Incorporate changes for Appendix B from 4/14/2021
0.5	R17	Bill Martin	27 April 2021	Accepted changes and removed resolved comments
0.5	R18	Bill Martin	27 May 2021	Added Configuration flow chart per proposal Updated Open CL CSS per proposal Updated figures in Usage section Updated Usage definition to make consistent Moved open comments to "editor's notes"
0.5	R19	Bill Martin	3 June 2021	Incorporated modifications to B.2
0.8	R0	Arnold Jones	9 June 2021	Formatted as a Working Draft for release outside the TWG per TWG approval ballot.
0.8	R1	Bill Martin	31 March 2022	Added initial security considerations Fixed must/shall wording throughout (3 instances) Added CSEE definition
0.8	R2	Bill Martin	1 June 2022	Removed track changes from previous revision Added modifications from SNIA-CS-Arch-Prog-Model-Theory-of-op-mods-06-01-2022.
0.8	R3	Bill Martin	21 June 2022	Removed previous change tracking Incorporated all RFC comments Resolved most RFC comments
0.8	R4	Bill Martin	22 June 2022	Removed previous change tracking Incorporated changes requested at TWG meeting All comments marked "done"

0.8	R5	Bill Martin	23 June 2022	Clean version – all comments removed
0.9		Bill Martin		Changed header from 0.8R5 to 0.9 for member vote and public review

Table of Contents

FOREWORD.....	13
1.1 SCOPE	13
2 REFERENCES	15
3 DEFINITIONS, ABBREVIATIONS, AND CONVENTIONS	16
3.1 DEFINITIONS	16
3.2 KEYWORDS.....	17
3.3 ABBREVIATIONS	18
4 THEORY OF OPERATION	19
4.1 OVERVIEW.....	19
4.2 DISCOVERY	22
4.3 CONFIGURATION	24
4.4 SECURITY.....	24
4.5 CSF USAGE	28
5 EXAMPLE COMPUTATIONAL STORAGE FUNCTIONS	32
5.1 COMPRESSION CSF	32
5.2 DATABASE FILTER CSF	32
5.3 ENCRYPTION CSF	32
5.4 ERASURE CODING CSF	32
5.5 REGEx CSF	33
5.6 SCATTER-GATHER CSF	33
5.7 PIPELINE CSF	33
5.8 VIDEO COMPRESSION CSF	33
5.9 HASH/CRC CSF.....	34
5.10 DATA DEDUPLICATION CSF	34
5.11 LARGE DATA SET CSFs	34

6	EXAMPLE COMPUTATIONAL STORAGE EXECUTION ENVIRONMENT	35
6.1	OPERATING SYSTEM CSEE	35
6.2	CONTAINER PLATFORM CSEE	35
6.3	CONTAINER CSEE	35
6.4	eBPF CSEE	35
6.5	FPGA BITSTREAM CSEE	35
ANNEX A.	(INFORMATIVE) ILLUSTRATIVE EXAMPLES	36
A.1	CSFs ON A LARGE MULTI-DEVICE DATASET USING CEPH.....	36
A.1.1	INTRODUCTION.....	36
A.1.2	THEORY OF OPERATION.....	38
A.1.3	DISCOVERY	38
A.1.4	CONFIGURATION	39
A.1.5	USAGE	39
A.1.6	EXAMPLE APPLICATION DEPLOYMENT	40
A.2	USING A CONTAINERIZED APPLICATION WITHIN LINUX (CSEE WITH INCLUDED CSF).....	42
A.2.1	THEORY OF OPERATION.....	42
A.2.2	DISCOVERY	42
A.2.3	CONFIGURATION	43
A.2.4	USAGE	44
A.2.5	EXAMPLE OF APPLICATION DEPLOYMENT	45
A.3	DATA DEDUPLICATION CSF	47
A.3.1	OVERVIEW.....	47
A.3.2	THEORY OF OPERATION.....	49
A.3.3	DISCOVERY	50
A.3.4	CONFIGURATION	50

A.3.5 OPERATION	52
A.3.6 MONITORING.....	52
A.3.7 CSF DATA DEDUPLICATION EXAMPLE.....	52
A.4 A COMPUTATIONAL STORAGE FUNCTION ON A NVM EXPRESS AND OPENCL BASED COMPUTATIONAL STORAGE DRIVE ILLUSTRATIVE EXAMPLE	58
A.4.1 CSEE EXAMPLE	58
A.4.2 COMPUTATIONAL STORAGE DRIVE.....	59
A.4.3 THEORY OF OPERATION.....	60
A.4.4 DISCOVERY	60
A.4.4.1 NVME FUNCTION DISCOVERY	60
A.4.4.2 OPENCL CSP FUNCTION DISCOVERY	61
A.4.5 CONFIGURATION – EXPLICIT MODE	62
A.4.6 USAGE – STORAGE DIRECT	62
A.4.7 USAGE – EXPLICIT MODE COMPUTATIONAL STORAGE	64
A.5 DATA COMPRESSION CSF EXAMPLE	66
A.5.1 OVERVIEW.....	66
A.5.2 THEORY OF OPERATION.....	66
A.5.3 DISCOVERY	66
A.5.4 CONFIGURATION	67
A.5.5 MONITORING.....	67
A.6 DATA FILTER CSF EXAMPLE.....	68
A.6.1 OVERVIEW.....	68
A.6.2 THEORY OF OPERATION.....	68
A.6.3 DISCOVERY	69
A.6.4 CONFIGURATION	69

A.6.5 OPERATION 69

A.6.6 MONITORING..... 70

Table of Figures

Figure 4.1– An Architectural view of Computational Storage.....	20
Figure 4.2 – CSF Discovery and Configuration Flowchart.....	23
Figure 4.3 - Direct Usage Model.....	29
Figure 4.4 – Indirect Usage Interactions	30
Figure B.1.1 – Ceph, a scale-out storage system.....	37
Figure B.1.2 – Ceph CSA	38
Figure B.1.3 – Ceph CSDs	38
Figure B.1.4 – Skyhook Example 1	40
Figure B.1.5 – Skyhook Command Execution	41
Figure B.2.6 – CSx discovery process.....	43
Figure B.2.7 – CSx Configuration process.....	44
Figure B.2.8 – CSx usage.....	45
Figure B.3.1 – Simplified Data Deduplication Process	47
Figure B.3. 2 Post Process Data Deduplication.....	48
Figure B.3. 3 Inline Data Deduplication	49
Figure B.3. 4 Configuration Parameters	51
Figure B.3. 5 Volume Creation	53
Figure B.3. 6 Enable Deduplication	53
Figure B.3. 7 Retrieve Status.....	54
Figure B.3. 8 Disable Deduplication Scheduling.....	54
Figure B.3. 9 Copy files	55
Figure B.3. 10 Verify space utilized	55
Figure B.3. 11 Schedule deduplication	56
Figure B.3. 12 Monitor Progress.....	56
Figure B.3. 13 verify space savings.....	57
Figure B.4.1 CS architectural diagram adapted for this illustrative example.....	59
Figure B.6. 1 An example to illustrate the function of data filter FCSS	68

FOREWORD

The SNIA Computational Storage Technical Working Group was formed to establish architectures and software computation in its many forms to be more tightly coupled with storage, at both the system and drive level. An architecture model and a programming model are necessary to allow vendor-neutral, interoperable implementations of this industry architecture.

This SNIA specification outlines the architectural models that are defined to be Computational Storage. As this specification is developed, requirements in interface standards and specific APIs may be proposed as separate documents and developed in the appropriate organizations.

1.1 Scope

This specification focuses on defining the capabilities and actions that are able to be implemented across the interface between Computational Storage devices (CSxes) (e.g. Computational Storage Processors, Computational Storage Drives and Computational Storage Arrays) and either Host Agents or other CSxes.

The actions mentioned above are associated with several aspects of a CSx:

- **Management.** Actions that allow Host Agent(s), based on security policies, to perform:
 - **Discovery.** Mechanisms to identify and determine the capabilities and Computational Storage Resources (CSR).
 - **Configuration.** Programming parameters for initialization, operation, and/or resource allocation
- **Security.** Considerations for security related to CSxes
- **Usage.** Allows a Host Agent or CSx to offload Computational Storage tasks to a CSx, including providing the target CSx with information about data locality both local to the CSx or resident on one or more non-local locations.

This specification makes no assumptions about the physical nature of the interface between the Host Agent and CSx(s). This specification and the actions associated with it will be implemented across a range of different physical interfaces. This specification also makes no assumptions about the storage protocols used by Host Agents and CSx(s).

The following storage protocols between the Host Agent and the CSx may be supported:

- **Logical Block Address.** Data is grouped into fixed-size logical units and operations are atomic at that unit size. Data is indexed via a numerical index into the Logical Block Address.
- **Key-Value.** Data is not fixed-size and is indexed by a key.
- **Persistent Memory.** Byte addressable non-volatile memory.

This specification defines actions for passing data through multiple Computational Storage Functions that may or may not reside on a single CSx. Additionally, it defines actions for requesting multiple Computational Storage Functions to perform a set of tasks.

2 References

The following referenced documents are indispensable for the application of this document.

SNIA Computational Storage API	Computational Storage API v0.5 rev 0 , available from https://www.snia.org/tech_activities/publicreview
NVMe® 2.0	NVM Express Base Specification 2.0, Approved standard, available from http://nvmexpress.org
FIPS 140-3	Federal Information Processing Standards Publication 140-3 (FIPS PUB 140-3) Security Requirements for Cryptographic Modules, March 2019, https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.140-3.pdf
ISO/IEC 19790	ISO/IEC 19790:2012(E), Information technology — Security techniques — Security requirements for cryptographic modules, August 2012, https://www.iso.org/standard/52906.html
ISO/IEC 24759	ISO/IEC 24759:2017(E), Information technology — Security techniques — Test requirements for cryptographic modules, March 2017, https://www.iso.org/standard/72515.html
ISO/IEC 15408 (multi-part standard)	ISO/IEC 15408-1:2009, Information technology — Security techniques — Evaluation criteria for IT security — Part 1: Introduction and general model, December 2009, https://www.iso.org/standard/50341.html ISO/IEC 15408-2:2008, Information technology — Security techniques — Evaluation criteria for IT security — Part 2: Security functional components, August 2008, https://www.iso.org/standard/46414.html ISO/IEC 15408-3:2008, Information technology — Security techniques — Evaluation criteria for IT security — Part 3: Security assurance components, August 2008, https://www.iso.org/standard/46413.html

3 Definitions, abbreviations, and conventions

For the purposes of this document, the following definitions and abbreviations apply.

3.1 Definitions

3.1.1 Allocated Function Data Memory

Function Data Memory (FDM) that is allocated for a particular instance of a CSF

3.1.2 Computational Storage

architectures that provide computation coupled to storage, offloading host processing or reducing data movement.

3.1.3 Computational Storage Array (CSA)

storage array that contains one or more CSEs.

3.1.4 Computational Storage Device (CSx)

Computational Storage Drive, Computational Storage Processor, or Computational Storage Array.

3.1.5 Computational Storage Drive (CSD)

storage element that contains one or more CSEs and persistent data storage.

3.1.6 Computational Storage Engine (CSE)

component that is able to execute one or more CSFs

note 1 to entry Examples are: CPU, FPGA.

3.1.7 Computational Storage Engine Environment (CSEE)

operating environment for a CSE

Note 1 to entry Examples are: Operating System, Container Platform, eBPF, and FPGA Bitstream.

3.1.8 Computational Storage Function (CSF)

a set of specific operations that may be configured and executed by a CSE.

Note 1 to entry Examples are: compression, RAID, erasure coding, regular expression, encryption.

3.1.9 Computational Storage Processor (CSP)

component that contains one or more CSEs for an associated storage system without providing persistent data storage

3.1.10 Computational Storage Resource (CSR)

resources available in a CSx necessary for that CSx to perform computation

Note 1 to entry Examples are: FDM, Resource Repository, CPU, memory, FPGA resources)

3.1.11 Function Data Memory (FDM)

Device memory used for storing data that is used by the Computational Storage Functions (CSFs) and is composed of allocated and unallocated Function Data Memory

3.1.12 platform firmware

the collection of all device firmware on a platform

3.2 Keywords

In the remainder of the specification, the following keywords are used to indicate text related to compliance:

3.2.1 mandatory

a keyword indicating an item that is required to conform to the behavior defined in this standard

3.2.2 may

a keyword that indicates flexibility of choice with no implied preference; “may” is equivalent to “may or may not”

3.2.3 may not

keywords that indicate flexibility of choice with no implied preference; “may not” is equivalent to “may or may not”

3.2.4 need not

keywords indicating a feature that is not required to be implemented; “need not” is equivalent to “is not required to”

3.2.5 optional

a keyword that describes features that are not required to be implemented by this standard; however, if any optional feature defined in this standard is implemented, then it shall be implemented as defined in this standard

3.2.6 shall

a keyword indicating a mandatory requirement; designers are required to implement all such mandatory requirements to ensure interoperability with other products that conform to this standard

3.2.7 should

a keyword indicating flexibility of choice with a strongly preferred alternative

3.3 Abbreviations

AFDM Allocated Function Data Memory

CSA Computational Storage Array

CSD Computational Storage Drive

CSE Computational Storage Engine

CSEE Computational Storage Engine Environment

CSF Computational Storage Function

CSP Computational Storage Processor

CSR Computational Storage Resources

CSx Computational Storage devices

FDM Function Data Memory

SSD Solid State Disk

4 Theory of Operation

4.1 Overview

This section describes the theory of operations for Computational Storage Devices (CSxes), Computational Storage Resources (CSRs), Computational Storage Engines (CSEs), Computational Storage Engine Environments (CSEEs), and Computational Storage Functions (CSFs).

Computational Storage architectures enable improvements in application performance and/or infrastructure efficiency through the integration of compute resources (outside of the traditional compute & memory architecture) either directly with storage or between the host and the storage. The goal of these architectures is to enable parallel computation and/or to alleviate constraints on existing compute, memory, storage, and I/O.

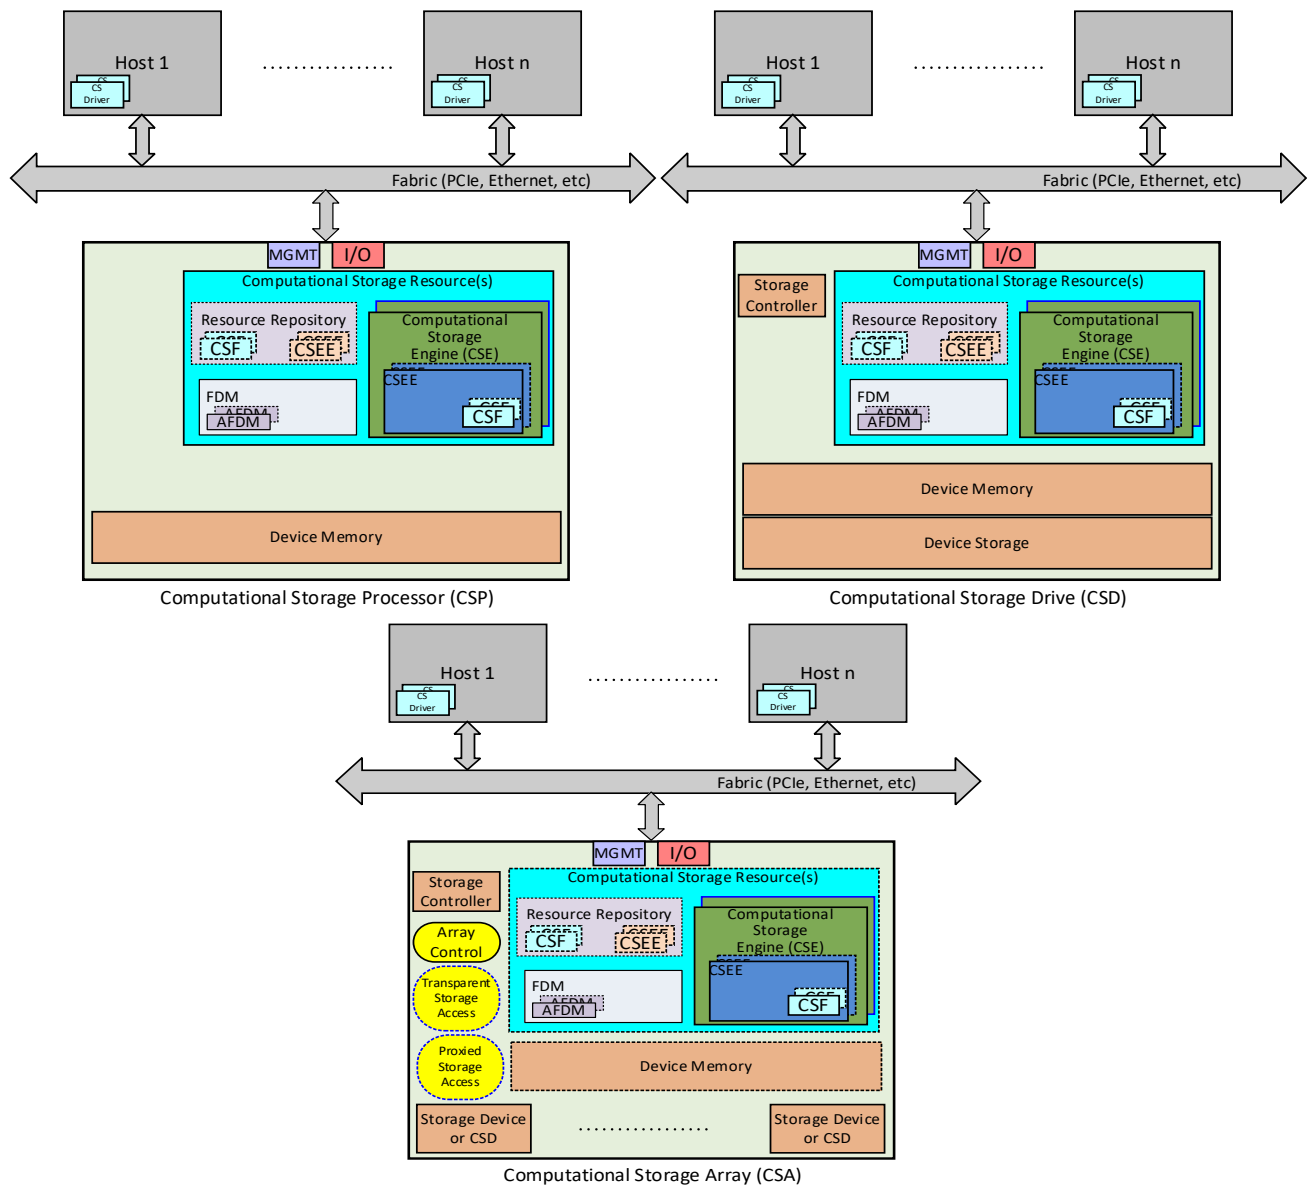


Figure 4.1– An Architectural view of Computational Storage

An illustrative example of Computational Storage devices (CSxes) is shown in Figure 4.1. A CSx consists of the following components:

- Computational Storage Resources (CSR) which contain:
 - A Resource Repository where the following may be stored:
 - Computational Storage Functions (CSFs); and
 - Computational Storage Engine Environments (CSEEs);
 - Function Data Memory (FDM) which may be partitioned into Allocated Function Data Memory (AFDM); and
 - One or more Computational Storage Engines (CSEs);
- A Storage Controller for CSD or an Array Controller for CSA;

- Device Memory; and
- Device Storage for CSD and CSA.

Computational Storage Resources (CSRs) are the resources available in a CSx necessary for that CSx to store and execute a CSF.

A Computational Storage Engine (CSE) is a CSR that is able to be programmed to provide one or more CSFs.

A Computational Storage Engine Environment (CSEE) is an operating environment for the CSE.

A Computational Storage Function (CSF) is a set of specific operations that may be configured and may be executed by a CSE in a CSEE.

Activation is the process of associating a CSEE with a CSE or associating a CSF with a CSEE. As part of activation of a CSEE, any resources that are necessary for that CSEE to be used on the CSE are assigned. As part of activation of a CSF, any resources that are necessary for that CSF to be used on the CSEE are assigned to the. When a CSEE association with a CSE or a CSF association with a CSEE is no longer required, the CSEE or CSF may be deactivated. This deactivation process releases any assigned resources.

A CSE is required to have a CSEE activated to be able to have a CSF activated. A CSE has FDM associated with it. A CSE is able to have one or more CSEEs and one or more CSFs activated at the time of manufacture that are usable by the host via management and I/O interfaces, or it is able to have one or more CSEEs and one or more CSFs downloaded by the host and activated. A CSE may have CSFs that have been programmed at the time of manufacture that are not changeable (i.e., not stored in the Resource Repository) (e.g., compression, RAID, erasure coding, regular expression, encryption). CSFs that are stored in the Resource Repository may be activated in a CSEE in a CSE.

A CSEE may be pre-installed or downloaded by the host. A downloaded CSEE or pre-installed CSEE is required to be activated for use. A CSEE may support the ability to have additional CSEEs activated within it. A CSEE may have a CSF embedded within the CSEE. That CSF may be implicitly activated when the CSEE is activated on a CSE.

A CSF is required to be activated on a CSE to be used. The CSF performs only the defined operations (e.g., a specific eBPF program or compression) that are reported by the CSx (i.e., the underlying operation is not changeable).

Function Data Memory (FDM) is device memory that is available for CSFs to use for data that is used or generated as part of the operation of a CSF. Allocated Function Data Memory (AFDM) is a portion of FDM that is allocated for one or more specific instances of a CSF operation. Any specific instance of a CSF operation shall only be allowed to access the AFDM allocated for that instance of a CSF operation. AFDM may be explicitly deallocated or may be deallocated on a power cycle or reset condition. As part of deallocating AFDM, the physical memory that was allocated to that AFDM should be cleared in order to prevent a subsequent instance of a CSF operation from accessing user data that is in that physical memory. On a reset, power cycle, or

sanitize operation that deallocates AFDM, all FDM should be cleared to prevent any instance of a CSF operation from accessing user data that was not specifically stored in AFDM for that instance of a CSF operation.

The Resource Repository is a region of memory and/or storage located within the CSx that contains zero or more images for CSFs and CSEEs that are available for activation. These CSFs and CSEEs are required to be activated in the CSE in order to be utilized.

A Computational Storage Processor (CSP) is a component that is able to execute one or more CSFs for an associated storage system without providing persistent data storage. The CSP contains CSRs and Device Memory. The mechanism by which the CSP is associated with the storage system is implementation specific.

A Computational Storage Drive (CSD) is a component that is able to execute one or more CSFs and provides persistent data storage. The CSD contains a Storage Controller, CSRs, Device Memory, and persistent data storage.

A CSD may function as a standard Storage Drive, with existing host interfaces and drive functions. As such, the system is able to have a storage controller with associated storage memory, along with storage addressable by the host through standard management and I/O interfaces.

A Computational Storage Array (CSA) is a storage array that is able to execute one or more CSFs. As a storage array, a CSA contains control software, which provides virtualization to storage services, storage devices, and CSRs for the purpose of aggregating, hiding complexity or adding new capabilities to lower level storage resources. The CSRs in the CSA may be centrally located or distributed across CSDs/CSPs within the array.

4.2 Discovery

4.2.1 CSx Discovery Overview

Discovery of CSxes is fabric dependent and is outside of the scope of this architecture.

4.2.2 CSR Discovery Overview

Once a CSx is discovered, to utilize Computational Storage Resources (CSRs), the characteristics of that CSx needs to be discovered. This involves a CSR discovery process for each discovered CSx. The CSR discovery process discovers all resources available including CSEs, CSEEs, CSFs, and FDM.

Discovery of a CSE includes information of any activated CSEEs and any activated CSFs in those CSEEs.

CSEEs in the Resource Repository may be discovered and information about any CSFs pre-activated in those CSEEs is returned. CSEEs in the Resource Repository are required to be activated in order to be used.

CSFs in the Resource Repository may be discovered. CSFs in the Resource Repository are required to be activated in order to be used.

Section 4.2.3 shows an example discovery flow. The specifics of a CSR discovery process are defined in API specifications (e.g., SNIA Computational Storage API).

4.2.3 CSF Discovery and Configuration Example

Figure 4.2 shows an example flowchart of discovery and configuration of a CSF. This example assumes that each of the actions can be completed and that there are no errors. This is only one example of how configuration is able to be completed.

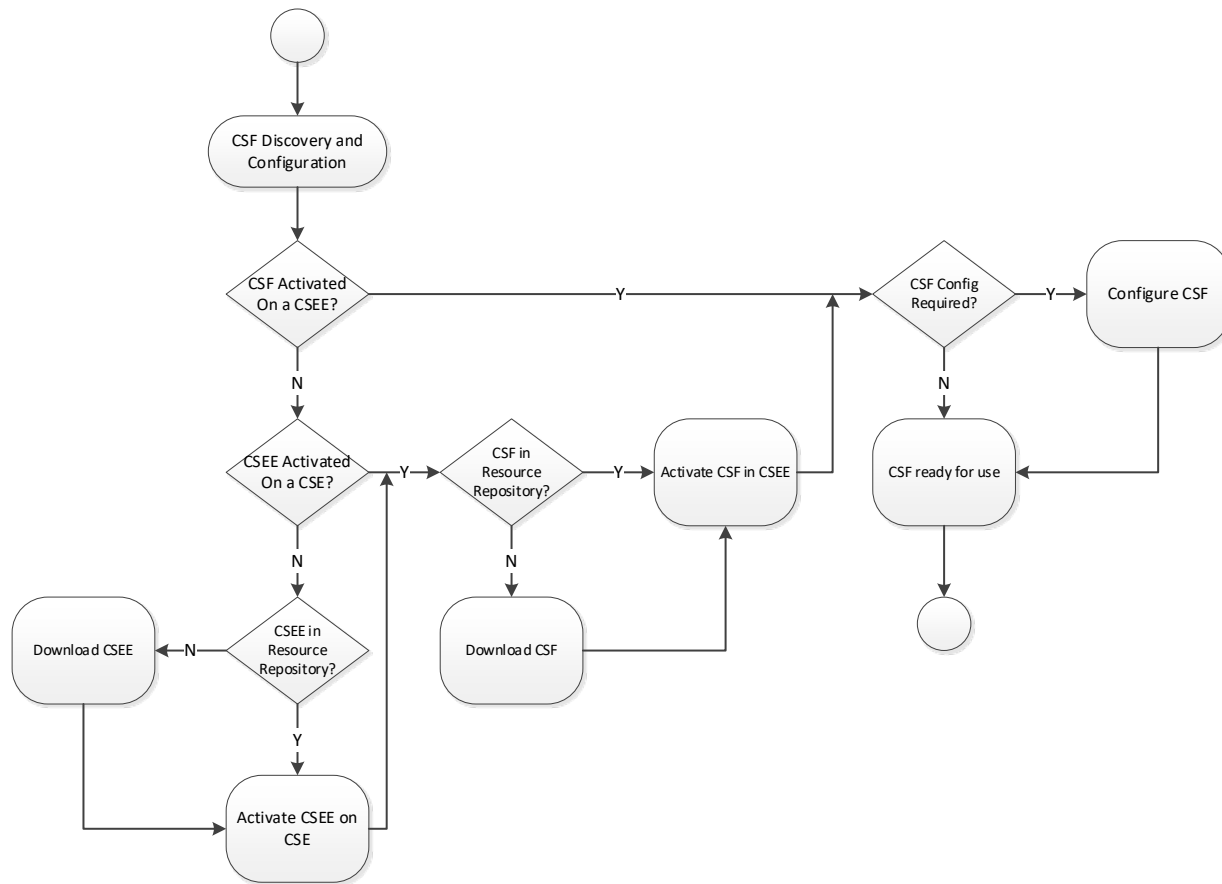


Figure 4.2 – CSF Discovery and Configuration Flowchart

The flow of the discovery and configuration process is a number of steps to **determine what is already activated in the CSX**. For a CSF or CSEE that is not already activated there is a discovery if the desired CSF or CSEE exists in the Resource Repository. If a desired CSF or

CSEE is not in the Resource Repository then it has to be downloaded to the Resource Repository.

For a desired CSEE that is not activated, that CSEE is required to be activated in a CSE. After the desired CSEE is activated and the desired CSF is available in the Resource Repository, that CSF is activated in the CSEE.

For a desired CSF that is activated in a CSEE, that CSF is configured, if there are static configurations that are required for all executions of that CSF. Once the CSF is configured it is available for an application to execute.

4.3 Configuration

4.3.1 CSE Configuration Overview

A CSE may be configured to prepare it for use. One aspect of CSE configuration is activation of one or more CSEEs. A CSEE may be activated in the CSE at time of manufacture and therefore not be required to be activated as part of configuration. The specifics of a CSE configuration process are defined in API specifications (e.g., SNIA Computational Storage API).

4.3.2 CSEE Configuration Overview

A CSEE may be configured to prepare it for use. A CSEE is required to be activated in order to be used by a CSE. One aspect of CSEE configuration is activation of one or more CSFs. A CSF may be pre-activated in the CSEE and therefore not be required to be activated as part of configuration. The specifics of a CSEE configuration process are defined in API specifications (e.g., SNIA Computational Storage API).

4.3.3 CSF Configuration Overview

A Computational Storage Function may be configured to prepare it for use. A CSF is required to be activated in order to be used by a CSEE. The specifics of a CSF configuration process are defined in API specifications (e.g., SNIA Computational Storage API).

This process may be done once for the CSF, prior to any specific invocation of the CSF, or as parameters associated with the invocation of a CSF.

4.4 Security

4.4.1 General

Security requirements for computational storage vary significantly (e.g., depending on the environment, interconnectivity, and sensitivity of data). As such, security is presented in this document as considerations that may be used to help determine the security that is appropriate to the risks. Some of the considerations are written such that specific requirements are identified for certain elements of security (e.g., a decision to use encryption results in specific requirements such as security strength, key management, and others).

The security considerations have been written with the following assumptions:

- a) the environment consists of a single physical host or virtual host with one or more CSxes
- b) the host is responsible for the security of the ecosystem that the CSxes operate within
- c) CSx security requirements are comparable to the security requirements common to SSDs/HDDs

It is important to note that multi-host environments as well as situations where the CSxes need to participate in the protection of data result in significantly more complex security considerations and requirements.

4.4.2 Privileged Access and Operations

As stated in the assumptions in 4.4.1, the host is responsible for much of the security in a basic computational storage configuration. Much of this security involves privileged operations that are performed/executed by individuals (systems administrators or other privileged users) and entities (e.g., system applications) that have elevated privileges that are beyond normal users and entities. While the security associated with these operations and accesses are out of scope for this document, this security is critical to protecting data, resources, configuration, and state.

4.4.3 CSx Security Considerations

The security considerations identified in this sub-section are those considered relevant, given the assumptions stated in 4.4.1. **Security is anticipated to be an important element of most computational storage implementations.**

A rogue or broken CSF could consume all of the resources that other CSFs may need to operate. Therefore, mitigation mechanisms (e.g. verification of CSFs and sandboxing of CSFs) should be considered.

Unless other steps are taken to prevent it, the associated storage for the CSx is consumable by any and all CSFs.

4.4.3.1 CSx Sanitization

As part of any Sanitize operation:

- a) all instances of CSF operations should be terminated
- b) all activated CSEEs and CSFs should be deactivated; and
- c) all memory allocations associated with FDM should be removed and associated FDM cleared.

4.4.3.2 CSx Data at-rest Encryption

Data at-rest encryption is a commonly required feature of storage devices. Therefore, a CSx should implement the same data at-rest encryption as would be implemented on any storage device in a similar application.

A CSx may include the capability to encrypt data prior to recording the resulting ciphertext on storage media and decrypt ciphertext that has been recorded on storage media. When data at-rest encryption is implemented, the following should be provided:

- a) strong symmetric encryption that provides a minimum of 128-bits of security strength to protect data (e.g., selection of encryption algorithms and modes of operations suitable for storage to be protected);
- b) cryptographic keys that are only used for one purpose (e.g., do not use key-encrypting keys (i.e., key wrapping keys) to encrypt data or use data encrypting keys to encrypt other keys);
- c) key management functionality (see 4.4.3.3) necessary for the data at-rest encryption;
- d) capability to rekey the data with a different data/media encryption key (DEK/MEK) (i.e., reading the data, decrypting it with the old key, encrypting the data with a new key, and writing the new ciphertext).

Additional elements of data at-rest encryption implementations may include:

- a) controls on the amount of data protected under a single key as well as within the established cryptoperiods;
- b) proof/verification of encryption (enabled v disabled);
- c) cryptographic modules used to protect sensitive or regulated data should be validated using recognized security criteria (e.g. ISO/IEC 19790, ISO/IEC 15408, and NIST FIPS 140-3); and
- d) archiving/escrowing the keys and keying material on key management servers.

The use of data at-rest encryption within a CSx has the following implications:

- a) import/export compliance issues may affect the sale, distribution, and use of the CSx in certain jurisdictions; which may require specific licensing; and
- b) data reduction technologies (e.g., compression and deduplication) are generally ineffective when applied to ciphertext.

4.4.3.3 CSx Key Management

Key management functionality is typically included in conjunction with data at-rest encryption (see 4.4.3.2) as opposed to a standalone capability.

A CSx may include key management capabilities to support encryption and decryption of data as well as cryptographic erase-based storage sanitization (see 4.4.3.4). When key management is implemented, the following should be provided:

- a) cryptographic services that provide a minimum of 128 bits of security strength;
- b) key generation with sufficient entropy (e.g., at least 256 bits of entropy input for AES-256) that uses the entire key space;
- c) secure distribution of the keys (e.g., authentication key or KEK);
- d) secure storage of keys and key material (e.g., with a hardware security module); and

- e) secure, secondary storage for key backup/recovery.

Providing Key Management within a CSx has the following implications:

- a) import/export compliance issues may affect the sale, distribution, and use of the CSx in certain jurisdictions; which may require specific licensing.

4.4.3.4 CSx Storage Sanitization

The controlled elimination of data in the form of storage sanitization is a commonly required feature of storage devices, therefore, it should be applied to CSxes. Failure to include storage sanitization may expose data to unauthorized access.

A CSx may include storage sanitization capabilities for controlled elimination of data. When storage sanitization is implemented, the following should be provided:

- a) storage sanitization (i.e., media-based or logical sanitization) using clear or purge methods;
- b) cryptographic erase (i.e., purge sanitization method in IEEE 2883) that ensures all copies of the encryption keys used to encrypt the target data are sanitized (see 4.4.3.2 data at-rest encryption and 4.4.3.3 key management); and
- c) validation of sanitization operation outcomes.

Additional elements of storage sanitization implementations may include:

- a) producing records (i.e., evidence) of sanitization operations that are able to serve as proof of sanitization; and
- b) sanitization performed in conjunction with autonomous data movement.

Providing storage sanitization within a CSx has the following implications:

- a) sanitization is not disrupted by firmware update, etc.

4.4.3.5 CSx Roots of Trust (RoT)

Roots of Trust (RoT) in the form of highly reliable hardware and software components that perform specific, critical security functions that provide a firm foundation from which to build security and trust are often included to support data at-rest encryption (see 4.4.3.2), key management (see 4.5.3.2), and attestation functionality.

If the CSx includes RoT and Chains of Trust (CoT), the following NIST SP 800-193, 4.1.1 requirements should be implemented:

- a) the security mechanisms are founded in Roots of Trust (RoT);
- b) if Chains of Trust (CoT) are used, a RoT serves as the anchor for the CoT;
- c) all RoTs and CoTs are either immutable or protected using mechanisms which ensure all RoTs and CoTs remain in a state of integrity; and

- d) all elements of the Chains of Trust for Update (NIST SP 800-193, 4.1.2), Detection (NIST SP 800-193, 4.1.3) and Recovery (NIST SP 800-193, 4.1.4) in non-volatile storage are implemented in platform firmware.

4.4.3.6 CSx Software Security

Software within a CSx may be in the following forms:

- a) computational Storage Engine Environment (CSEE); or
- b) computational Storage Functions (CSFs).

This software may be pre-installed or downloaded by the host.

A CSx may include a wide range of software security mechanisms, but the following should be provided:

- a) verification of the integrity of downloaded CSx software (e.g., use of checksums to detect errors);
- b) validation that the code is coming from a particular source and that the code has not been altered or compromised by a third party (e.g., use of code signing).

4.5 CSF Usage

4.5.1 CSF Usage Overview

Once configured, a host may use the CSF with:

- a) a direct usage model; or
- b) an indirect usage model.

In the direct usage model, the host sends a computation request that specifies a CSF to execute on data in the FDM. The data movement between host or storage and the FDM may be done outside of the operation of the CSF.

In the indirect usage model, the host sends a storage request to the Storage Controller. A CSF is executed on the data associated with a storage request based on:

- a) parameters in the storage request;
- b) the data locality; or
- c) the data characteristics (e.g., size).

For the indirect usage model that operates on data based on locality or characteristics, the Storage Controller is configured to associate a CSF with data locality or data characteristics prior to sending a storage request.

4.5.1.1 Direct CSF Usage Model

Figure 4.3 shows an example of the direct CSF usage model.

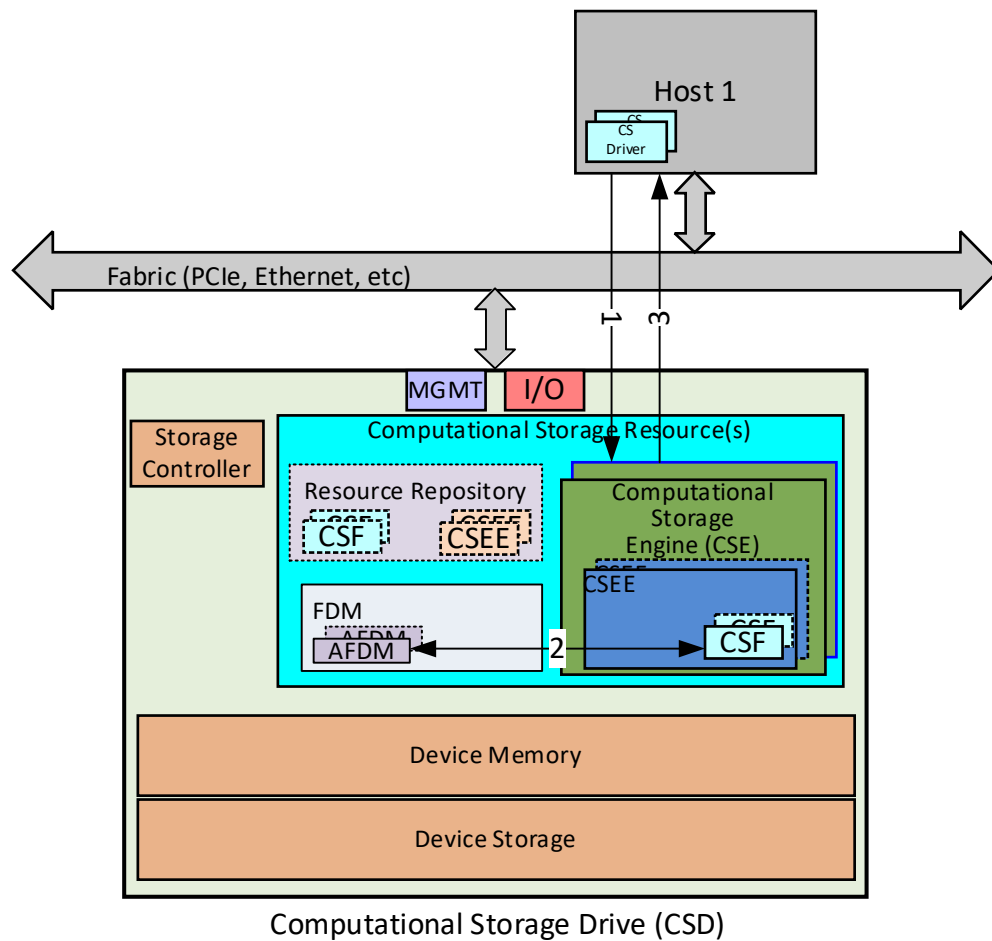


Figure 4.3 - Direct Usage Model

Figure 4.3 assumes that AFDM is allocated for the specific instance of the CSF and that data on which computation is to be performed is placed in that AFDM, prior to the request to the CSE, through some process that is not shown in this figure. The result data, if any, is placed in the AFDM and through some process, not shown in this example, moved from the AFDM to storage or the host. The steps shown in Figure 4.3 for a direct usage model are:

- (1) The host sends a command to invoke the CSF;
- (2) The CSE performs the requested computation on data that is in AFDM and places the result, if any, into AFDM; and
- (3) The CSE returns a response to the host.

4.5.1.2 Indirect CSF Usage Model

Figure 4.4 shows an example of the indirect CSF usage model.

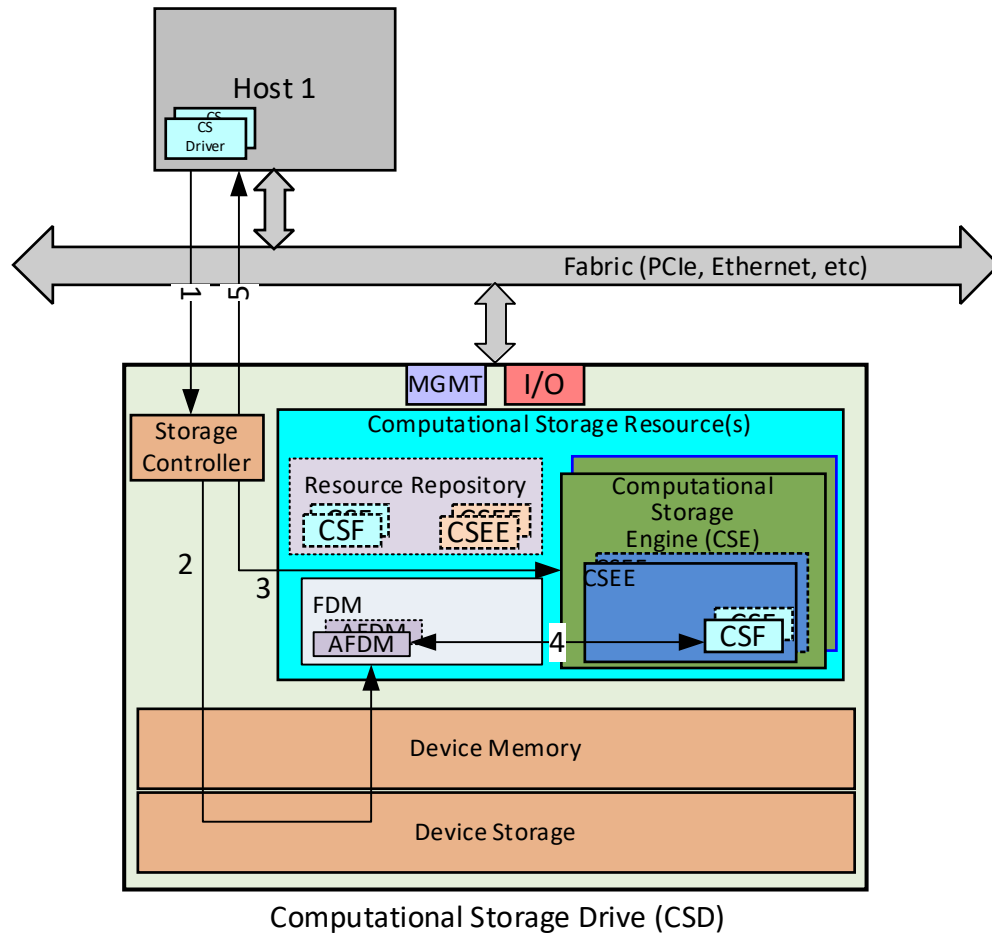


Figure 4.4 – Indirect Usage Interactions

Figure 4.4 assumes a read operation with computation on the data that is being read. The steps shown in Figure 4.4 to perform an indirect computation through the Storage Controller are:

- (1) The host configures the CSD to associate a specific CSF with reads that have specific characteristics;
- (2) The host sends a storage request to a Storage Controller where:
 - a. that storage request is associated with that target CSF; and
 - b. the storage controller determines what CSF is associated with the storage request;
- (3) The Storage Controller moves data from storage into the FDM;
- (4) The Storage Controller instructs the CSE to perform the indicated computation on the data in the FDM;
- (5) The CSE performs the computation on the data and places the result, if any, into the FDM; and
- (6) The Storage Controller returns the computation results, if any, from the FDM to the host.

4.5.2 CSF Execution

A CSF execution is specific to the type of CSF (e.g., for a compression CSF, a command may instruct the CSF to read from a given location in system memory, compress the data, and store the resulting data to a specified location in a storage device).

5 Example Computational Storage Functions

This section describes example Computational Storage Functions (CSFs) (see 4.1).

See section (See 4.2 and 4.3) for information about CSF discovery and configuration.

5.1 Compression CSF

A compression CSF reads data from a source location, compresses or decompresses the data, and writes the result to a destination location.

CSF configuration specifies the compression algorithm and associated parameters.

CSF command specifies the source address and length and the destination address and maximum lengths.

5.2 Database Filter CSF

A database filter CSF reads data from source location(s), performs a database projection (column selection) and filter (row selection) on the data according to projection and filter conditions, and writes the result(s) to destination location(s).

CSF configuration specifies the database format, table schema, selection and filter conditions, and associated parameters.

CSF command specifies the source address and length, and the destination addresses and lengths.

5.3 Encryption CSF

An encryption CSF reads data from a source location, encrypts or decrypts the data, and writes the result to a destination location.

CSF configuration specifies the encryption algorithm, keying information, and associated parameters.

CSF command specifies the source address and length and the destination address and length.

5.4 Erasure Coding CSF

An erasure coding CSF reads data from source location(s), performs a EC encode or decode on the data, and writes the result(s) to destination location(s).

CSF configuration specifies the EC algorithm and associated parameters.

CSF command specifies the source address and length and the destination addresses and lengths.

5.5 RegEx CSF

A regex CSF reads data from source location(s), performs a regular expression patterning matching or transformation on the data, and writes the result(s) to the destination location.

CSF configuration specifies the RegEx string(s) and associated parameters.

CSF command specifies the source address and length and the destination address and length.

5.6 Scatter-Gather CSF

A Scatter-Gather CSF reads data from set of source location(s) and writes the data to a set of destination location(s).

CSF configuration does not have any parameters.

CSF command specifies the source addresses and lengths and the destination addresses and lengths.

5.7 Pipeline CSF

A Pipeline CSF performs a series of operations on data according to a data flow specification, allowing different CSF commands to be combined together in a standardized way.

CSF configuration does not have any parameters.

CSF command specifies a collection of commands, their order and dependencies, and calculations defining the relationships of the addresses between commands.

5.8 Video Compression CSF

A video compression CSF reads data from a source location, compresses or decompresses the video, and writes the result to a destination location. In order to accommodate multiple parallel compressions, the video compression CSF may support a single compression stream or multiple compression stream

CSF configuration specifies the stream, compression algorithm and associated parameters.

CSF command specifies the stream, source address and length and the destination address and maximum lengths.

5.9 Hash/CRC CSF

A hash/CRC CSF reads data from a source location, calculates a hash or CRC value based on the source data, and writes the result to a destination location.

CSF configuration specifies the hashing/CRC algorithm and associated parameters.

CSF command specifies the source address and length and the destination address.

As an optional feature this CSF may calculate the hash/CRC value based on the source data and compare the hash/CRC result to a pre-calculated value supplied by the initiator. The CSS will notify the initiator whether the calculated value matches the supplied value.

5.10 Data Deduplication CSF

A data deduplication CSF reads data from source location(s), performs deduplication or duplication on the data, and writes the result(s) to the destination location(s). CSF configuration specifies the data deduplication algorithm and associated parameters.

CSF command specifies the source address and length and the destination address and maximum lengths.

5.11 Large Data Set CSFs

This example is for a large data set wherein the data is sharded as objects across a plurality of computational storage devices (CSxes) and these objects are further tagged as belonging to a named object class. The object class being defined as a set of methods that act on those named objects. The object class is the CSF. The object class subsystem is the CSE. There are CSEs defined and configured in each of the CSxes. The FDM, and AFDM is pulled from system memory CSF configuration includes the object class methods.

The CSF command specifies the objects names to be acted upon, the object class method to enact and other parameters for the object class model.

6 Example Computational Storage Execution Environment

This section describes example Computational Storage Execution Environments (CSEEs) (see 4.1).

See section (see 4.2 and 4.3) for information about CSEE discovery, configuration, and activation.

6.1 Operating System CSEE

An Operating System CSEE provides a specific operating system environment (e.g., Linux). The Operating System CSEE may contain one or more activated CSFs and may support the activation of one or more downloaded CSFs.

6.2 Container Platform CSEE

A Container Platform CSEE provides an environment to host one or more Container CSEEs.

In order to provide CSFs, it is necessary to have this type of CSEE configured with a Container CSEE.

6.3 Container CSEE

A Container CSEE provides a container environment. The Container CSEE may contain one or more activated CSFs and may support the activation of one or more downloaded CSFs.

6.4 eBPF CSEE

An extended Berkeley Packet Filter (eBPF) CSEE provides an environment for running eBPF programs. The eBPF CSEE may contain one or more activated eBPF CSFs and supports the activation of one or more downloaded eBPF CSFs.

6.5 FPGA Bitstream CSEE

A FPGA Bitstream CSEE provides an environment for an FPGA device. The FPGA Bitstream CSEE may contain one or more activated CSFs and may support the activation of one or more downloaded CSFs.

Annex A. (Informative) Illustrative Examples

A.1 CSFs on a Large Multi-Device Dataset using Ceph

A.1.1 Introduction

A large multi-device dataset is a dataset that:

1. may not fit into a single storage device;
2. may be large enough to require hundreds or thousands of devices; and
3. may require scalable performance by utilizing many storage devices.

Due to the size of these datasets, a single server may be insufficient to house all of the necessary storage devices. Consequently, these datasets may also span servers. This illustrative example uses TCP/IP as it provides scaling for a large number of devices.

A large dataset is to be sharded into chunks, that have semantic meaning to the application and are stored across a set of storage devices. To act on that data in the computational storage sense it is necessary to map the data shards to the devices where they are stored and then deliver a function to each of the devices. That function is then able to be executed in the device against each of the data set's shards stored in that device. This may be done simultaneously on thousands of devices.

There are many systems that enable the scaling of storage to thousands of devices. One such system used for this example is Ceph. Ceph allows many applications to jointly share shards of data called objects across potentially thousands of devices. Ceph is responsible for mapping the location of each of the objects across all the devices. Although intermediary servers called object storage daemons (OSDs) use local storage interconnects, the primary application interconnect is TCP/IP. Applications locate and interact with an object by a unique key that translates to a unique IP and TCP port address. Applications do not dictate this address but rather let Ceph manage the location of the object, abstracting the clients from the actual location.

Below is a diagram of Ceph showing client applications running in containers (App CT) using a variety of APIs (File, Block, S3) that are all implemented using the underlying Ceph RADOS API. This API permits the storing and retrieving of arbitrary sized objects as well as executing methods against objects.

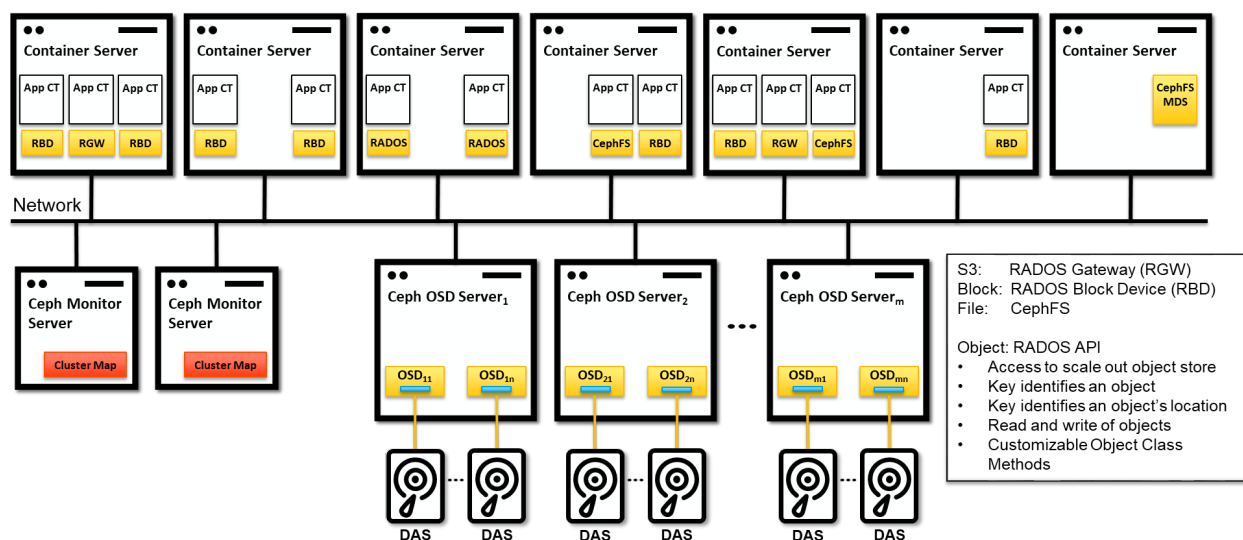


Figure B.1.1 – Ceph, a scale-out storage system

The Ceph OSD servers in the diagram are responsive to the application's object requests. Although a single OSD satisfies a single object request, a dataset may be sharded into many objects and those objects will be stored across all of the available OSDs. Since the application is abstracted away from an object's location, the application should also be abstracted away from the location of the execution environment. Furthermore, operating on a large dataset means engaging with many OSDs so abstracting the workload above the device level is also appropriate (i.e., an application should be able to have a distributed dataset while still operating on it as a single dataset).

Looking at the Computational Storage Architecture, Ceph could be viewed as the CSA and the OSDs could be viewed as Computational Storage Processors.

A CSE would be discovered and configured in the OSD.

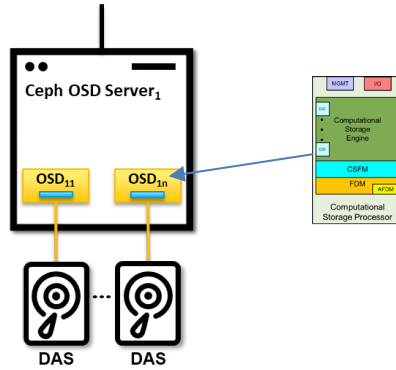


Figure B.1.2 – Ceph CSA

However, nothing in the software architecture prevents running a Ceph OSD as a single device server. In this case, the OSD server would be viewed as a Computational Storage Drive (CSD).




Figure B.1.3 – Ceph CSDs

A.1.2 Theory of Operation

In this section, we walk through the Large Multi-Device Dataset Ceph Computational Storage Theory of Operation (See Section 4) and apply it in a specific manner to this illustrative example. The three main phases of operation are:

- Discovery (see section 4.2)
- Configuration (see section 4.3)
- Usage (see section 4.4)

A.1.3 Discovery

Most of the device discovery in this example is handled by Ceph; however, discovery of the Ceph system and the available CSRs and CSFs needs to be done. Built into Ceph is the ability to run code in the OSDs against an object by means of an object class definition. This facility should be viewed as the CSE. An object class definition contains methods that can be executed within an OSD against an object. An object class should be viewed as a CSF. Both preloaded and downloaded CSFs are able to be created, so discovery of these object classes will be necessary.

A.1.4 Configuration

for this illustrative example, a Ceph system is required to be preinstalled and configured. The workloads are deployed by delivering the CSFs (i.e., object classes, if required) to each of the CSEs (i.e., OSD servers).

A.1.5 Usage

The application calls the CSF using the CSF parameters. Abstracted from the application, the CSF is translated into the Ceph object class. The CSF parameters are the object class method, the name(s) of the targeted objects, and any inbound parameters. The invocation of the CSF is handled by the Ceph RADOS `rados_exec()` function¹. Upon return any outbound parameters are returned including the calls status.

¹ There are several variants of the `rados_exec()` call, including versions that are associated with a read and write. There are also async versions of each call.

A.1.6 Example Application Deployment

Skyhook is an opensource middleware application that provides a Ceph backend to PostgreSQL. It enables PostgreSQL table stores by sharding tables on row boundaries, placing N rows of a table into a single object. If $N = 1000$ and an example table had a total of 100,000 rows, Skyhook would create 100 objects that would be distributed to potentially 100 OSDs/Devices in Ceph.

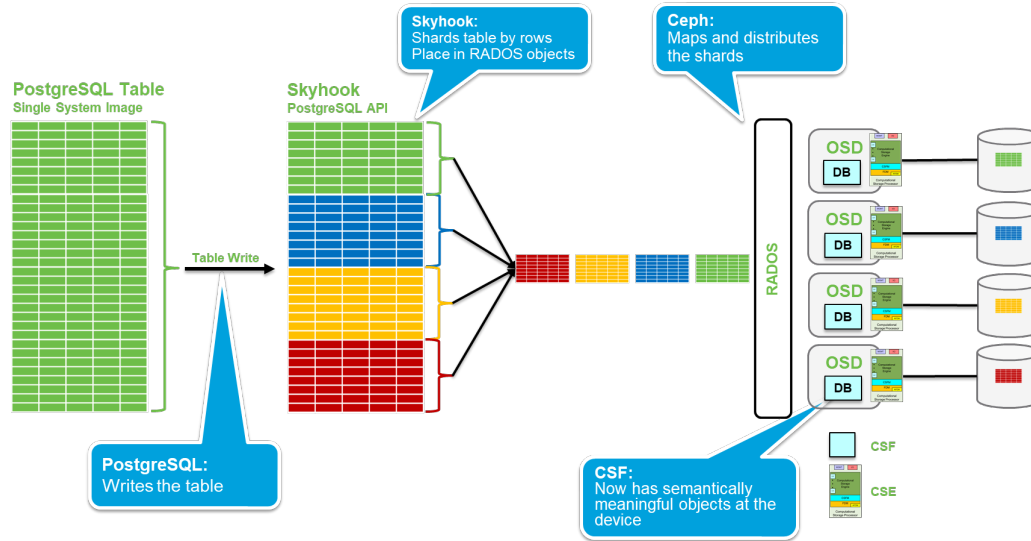


Figure B.1.4 – Skyhook Example 1

Skyhook has defined its own object class (CSF) on each Ceph OSD node (CSE). This object class implements methods that evaluate SQL queries. When PostgreSQL receives a user submitted SQL query, it submits the query to Skyhook. Skyhook then submits that query to each of the 100 objects of the table in parallel by calling the RADOS `exec()` function (CSF Command). This call (CSF command) passes the query to the SQL evaluation method, each object independently evaluates the query for its rows, returning the results back to skyhook. Skyhook then assembles all of the results into a single result that is then returned to PostgreSQL.

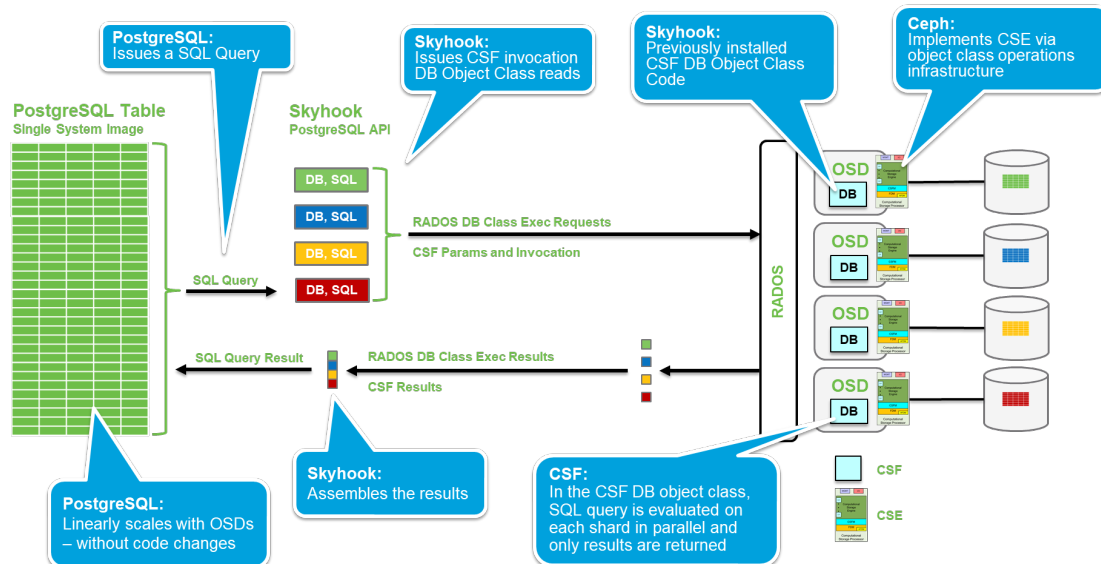


Figure B.1.5 – Skyhook Command Execution

Not only does this allow the DB to scale the performance and size independent of device capacities, but it also implements parallel execution.

A.2 Using a Containerized Application within Linux (CSEE with included CSF)

This illustrative example is of a Computational Storage Drive (CSD) based on the NVMe™ over PCIe® specification, that consists of one typical LBA based NVM storage device, multiple programmable applications processors co-located on the same controller capable of running Linux OS capable of running containerized application for searching data.

Specific assumptions for this illustrative example include:

- a) A server running a modern Linux kernel and user-space distribution;
- b) A single-ported single host system (i.e., no consideration for multi-port NVMe devices;
- c) No use of virtualization;
- d) A pre-activated CSEE (Running a Linux OS - CSEE1) is within the CSE;
- e) A CSEE is available to store in the Device Storage that is a Container (CSEE2) with a CSF1;
- f) CSF1 runs an Artificial Intelligence (AI) Application;
- g) No Peer-to-Peer capabilities or namespaces are used; and
- h) Leverages existing PCIe and NVMe methods around security.

A.2.1 Theory of Operation

In this section, we will walk through the Theory of Operation (see section 4) and apply it in a specific manner to this illustrative example. The three main phases of operation are:

- a) Discovery;
- b) Configuration; and
- c) Usage.

In this illustrative example, it is assumed that some updates to the existing NVMe Linux driver are made. The driver is required to be capable of recognizing and configuring a Computational Storage drive.

A.2.2 Discovery

NVMe/PCIe already has a very robust controller discovery and creation process. A PCIe device (i.e. an NVMe™ controller) is able to be discovered by the host at power-up time via PCIe bus enumeration or at run-time via a hot-plug event and bus rescan. A modern Linux operating system is able to detect PCIe devices via both the methods mentioned above, allowing discovery of a new NVMe controller to be done at any time in a running system.

Once an NVMe controller has been detected, the NVMe Computational Storage driver is able to be used to discover the capabilities of the controller via the NVMe Identify Admin command. The procedure is:

1. PCIe enumeration discovers the NVMe controller of the CSD;
2. The Computational Storage NVMe Linux driver binds to this PCIe device; and

3. The driver discovers the capabilities of the Computational Storage Drive:
 - a. An activated CSEE within the CSE that runs a Linux OS Environment (CSEE1).




Figure B.2.6 – CSx discovery process

A.2.3 Configuration

The CSEE1 is active within the CSE. As shown in Figure B.2.7, the following steps take place:

- A) the CSEE2 for the Docker Container is moved into the Device Storage of the CSD. An Admin command causes the application processor(s) to boot into a Linux Environment. The software that supports Docker is loaded automatically as part of Linux start-up process; and
- B) once Linux has booted, the CSEE2 for the Docker Environment is loaded with the CSF1.

The Computational Storage NVMe driver is notified so that the specific compute workload can be downloaded.

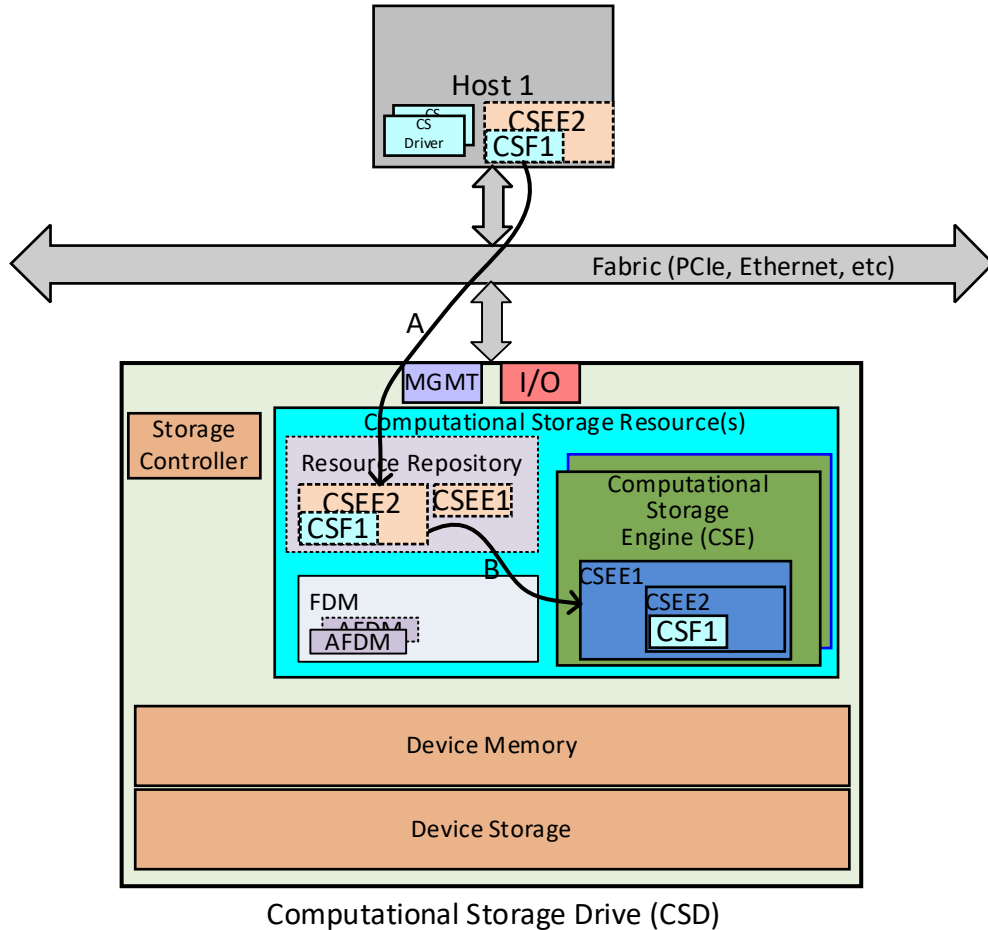


Figure B.2.7 – CSx Configuration process

A.2.4 Usage

Once the configuration process has completed and both CSEE1 and CSEE2 are activated and the CSF1 is ready to process data, the usage process steps (as shown in Figure B.2.8) can occur:

- 1) The CSF1 is told by the host to execute the function on data;
- 2) Input Data is pulled from Device Storage into Device Memory, in the Allocated Function Data Memory (AFDM) space of the Function Data Memory (FDM);
- 3) The data is operated on by the CSF1 and is then stored as Output Data in the Device Storage at the location specified by the function; and
- 4) The host can then access both the original Input Data and the new Output Data as required to complete the process steps.




Figure B.2.8 – CSx usage

A.2.5 Example of Application Deployment

The open source application called Openalpr (Open source Automatic License Plate Recognition) may be deployed using the Docker market place.

OpenALPR is an open source *Automatic License Plate Recognition* library written in C++ with bindings in C#, Java, Node.js, Go, and Python. The library analyzes images and video streams to identify license plates. The output is the text representation of any license plate characters. Openalpr provides a set of shared libraries but also makes use of a few other shared open source libraries.

In the Docker market place, there is a Docker image based on the “Vendor OS Choice” that contains the Openalpr shared libraries, its command-line utility application, and all the required shared libraries (e.g., OpenCV, python, and java)

In this specific example, the user would:

- 1) ssh from the host to the device using the default IP/username/password provided by the **device** vendor. (e.g., `user@server:~# ssh csd@ipaddress`)
- 2) Build a Docker image (i.e., `user@server:~# docker build -t openalpr https://github.com/openalpr/openalpr.git`)
- 3) Download test image (i.e., `user@server:~# wget http://plates.openalpr.com/h786poj.jpg`)
- 4) Run alpr on image (i.e., `user@server:~# docker run -it --rm -v $(pwd):/data:ro openalpr -c eu h786poj.jpg`)

The output of this example is:



```
user@server:~# /openalpr$ alpr ./h786poj.jpg
```

plate0: top 10 results -- Processing Time = 58.1879ms.

- PE3R2X confidence: 88.9371
- PE32X confidence: 78.1385
- PE3R2 confidence: 77.5444
- PE3R2Y confidence: 76.1448

A.3 Data Deduplication CSF

A.3.1 Overview

Data deduplication is a technique used to reduce the amount of data stored on a storage device. Compression results in the removal of repeating bytes or streams within a chunk or segment of storage, but data deduplication results in the removal of matching chunks or segments of storage.

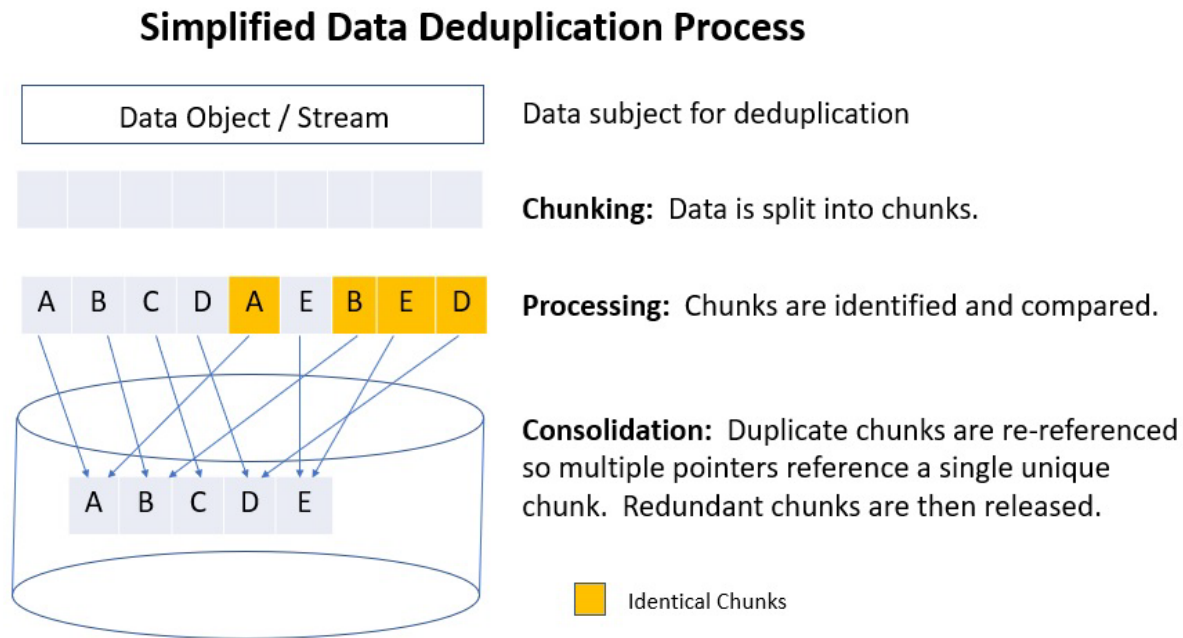



Figure B.3.1 – Simplified Data Deduplication Process

Figure B.3.1 describes a simplified data deduplication process where a given data object or stream is given to the deduplication process. The data object is split into chunks where the chunks can then be identified and compared. A location repository of pointers referencing the unique chunks is created and any duplicate chunks are released from the data object so that the resulting storage is reduced.

There are two generic ways of performing data deduplication on a device. The first way is post process data deduplication where the data that is being deduplicated already resides on the device and is processed at a scheduled time and duration. The second way is inline data deduplication where the data that is being deduplicated is immediately processed for deduplication so that only unique data is stored on the device.

Post Process Data Deduplication



1. Data object or stream is saved directly to the device.
2. The data is chunked with metadata and a repository is created to identify each chunk.
3. Chunks are compared. The chunk repository and chunk metadata are updated if chunks are identical.
4. The identical chunks are freed so that there is only a single instance of a chunk that was originally stored on the device.
5. The chunks are then consolidated within the device.

Identical Chunks

Figure B.3. 2 Post Process Data Deduplication

Figure B.3. 2 describes post process data deduplication. This process does not interfere with the initial ingestion of the data object or stream and allows the system to schedule a time and duration for the deduplication process. Once the data object or stream is stored, the data is chunked where the chunk metadata describes the size, references, and location of the data. A repository is also created to map the location of each chunk with respect to the way the data was initially stored. The chunks are compared and if chunks are identical, the chunk metadata of the first instance of the chunk is updated with a reference count. The repository pointer for the identical chunk is also updated to the first instance of the chunk. Once the data object or stream has had all the chunks compared, the identical chunks are freed so that only a single instance of the chunk is stored on the device. Finally, the chunks are consolidated within the device.

Inline Data Deduplication

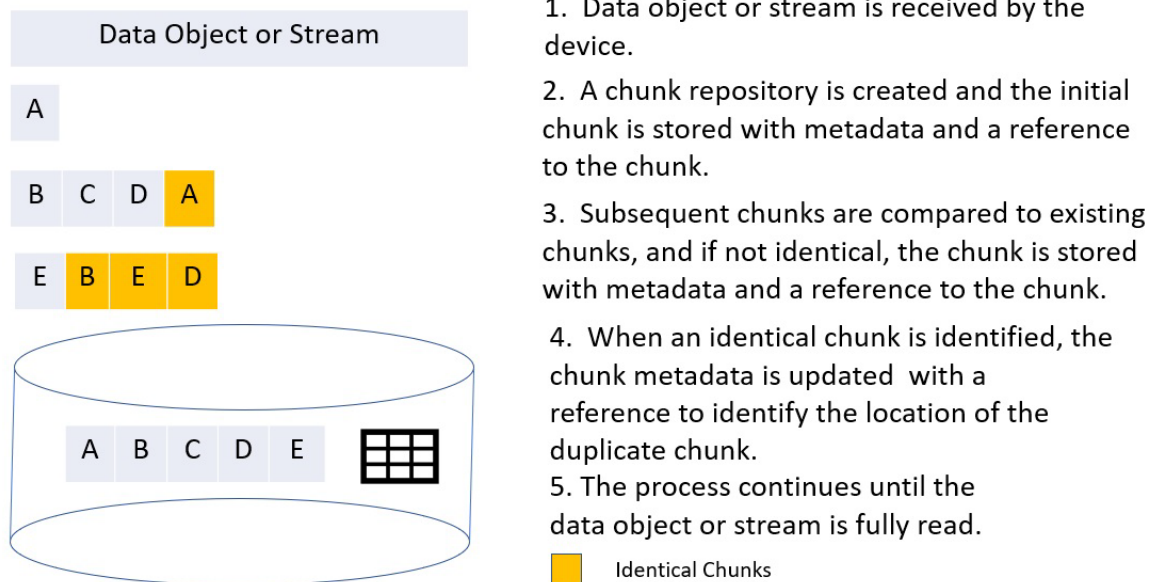


Figure B.3. 3 Inline Data Deduplication

Figure B.3.3 describes inline data deduplication. This process interferes with the ingestion of the initial data object or stream being saved to the device. As the data object or stream is being written to the device, chunks are compared and written only once if identical. This allows for less storage space to be used on the initial ingestion of the data object or stream. The first time the data is stored into the device, a chunk repository is created to describe the location of the chunks within the data object or stream. The initial chunk is stored with chunk metadata that describes size, location, and references to the chunk. Subsequent chunks are then compared with existing chunks, and if unique, it is stored as a unique chunk. If it's identical, the repository is updated to reference the single instance of the chunk and the number of references in the chunk metadata is updated. The process continues until the data object or stream is complete.

These descriptions of data deduplication are all simplified to explain generic ways to perform data deduplication. The CSx that will implement data deduplication will likely have more intricate proprietary ways of performing the data deduplication.

A.3.2 Theory of Operation

To successfully implement data deduplication as a CSF, a CSx must be able to communicate the ability to perform data deduplication within a CSE. Once the ability is determined, the CSF then needs to be successfully configured. The data object or data stream can then be processed by the CSE using the data deduplication CSF while allowing monitoring of the progress of that operation.

This illustrative example will attempt to provide the initial framework of the following steps needed to successfully allow for a CSx to perform data deduplication as a CSF.

A.3.3 Discovery

To determine if a CSx supports the Data deduplication CSF, the CSx needs to first be discovered as a CSx. If a CSx supports the Data deduplication CSF, then the CSx needs to also indicate if the Data deduplication CSF allows configuration. If the CSx allows for certain configuration parameters to the Data deduplication CSF, then the configurable parameters need to be shared. The following is a possible list of configurable parameters:

1. Supported Chunk Sizes – The acceptable size values to chunk the data object into comparable chunks. Typically, these are large, but they could also be variable.
2. Scheduling – The schedule and duration of post process data deduplication.
3. Failover – The action to take if data deduplication is interrupted (e.g., discarding, resuming from last good write, restarting).
4. Monitoring – The type of data to collect during data deduplication (e.g., the current data deduplication space savings, the I/O rate of the data deduplication operation, the size of the data processed, the size of the data remaining to be processed, the percentage complete).
5. Inline or Post Process Deduplication
6. Type of method to perform chunk comparison – Three common types are hashing, binary comparison and delta differencing.
7. Hashing Algorithm – Type of hashing algorithms allowed to identify unique chunks.
8. Data deduplication Analysis – Methods to determine the likely savings gained by performing data deduplication on a data object or stream.
9. Operational Interruption – Whether or not the data deduplication operation is able to be interrupted for purpose of either suspending, abandoning or resuming the request.

Note that discovery may also return the default values of the configurable parameters as well as capabilities like operational interruption that the Data deduplication CSF supports.

A.3.4 Configuration

Once a CSx that supports the Data deduplication CSF is discovered, the Data deduplication CSF can be configured if allowed. It's possible that the Data deduplication CSF will have default values that the user may not want to override, and configuration is not necessary. If the user wants to configure the Data deduplication CSF, then the parameters returned by discovery need to be set and sent to the Data deduplication CSF before the data deduplication operation is performed by the CSE.

Table 1

Element Name	Requirement	Description
--------------	-------------	-------------

DeduplicationSupport	Mandatory	Specifies whether Deduplication can be configured as a CSF within the CSE.
DeduplicationType	Optional	The type of deduplication to perform. This can be either inline or post process deduplication
DeduplicationSchedule	Conditional	Valid when post process deduplication is specified. This would be the time, frequency, and duration when the deduplication would be performed.
DeduplicationChunk	Optional	If set, the size of the chunks to perform deduplication comparisons.
DeduplicationFailover	Optional	If set, the action to take when data deduplication is interrupted. Possible actions are discarding, resume from last good write, restart, etc.
DeduplicationMonitoring	Optional	If set, type of data to collect during data deduplication. Possible types of data to collect are current data deduplication space savings, the I/O rate of the data deduplication operation, the size of the data processed, the size of the data remaining to be processed, the percentage complete, etc.
DeduplicationComparison	Optional	If set, type of method to perform chunk comparison. Three common types are hashing, binary comparison and delta differencing
DeduplicationHash	Optional	If set, type of hashing algorithm to identify unique chunks
DeduplicationSavings	Optional	If set, perform an analysis on data being received to determine the amount space saved.
DeduplicationOperation	Optional	If set, allow for the interruption of the data deduplication function by request with option to either suspend, abandoned or resume the data deduplication operation.

Figure B.3. 4 Configuration Parameters

Figure B.3.4 lists the type of parameters that could be configured for the Data deduplication CSF prior to sending a data object or stream to the CSx.

A.3.5 Operation

The data deduplication CSF will begin when the data object or stream is sent. The data object or stream may have already been processed by another CSF or be processed for input into another CSF. For example, data deduplication is inefficient on encrypted data, so the Encryption CSF should first decrypt the data before sending it to the data deduplication CSF. Additionally, once the data has been deduplicated, the data is able to then be sent to the compression CSF for additional storage savings.

The operation of the data deduplication CSF may be interrupted if allowed. Otherwise, the user will need to wait for the operation to complete or fail to determine the next course of action.

A.3.6 Monitoring

As the data object is being processed by the data deduplication CSF, the user can get status and statistics on the process. The following are possible status and statistics to monitor:

1. Current data deduplication space savings
2. The I/O rate of the data deduplication operation
3. Current amount of data processed by the data deduplication operation
4. Current amount of data remaining to be processed by the data deduplication operation
5. The percentage of completion of the data deduplication operation
6. Success or failure of the operation
7. Existing state of the operation such as paused, interrupted, resumed, etc.

Based on the status and statistics, the user can then determine if the operation needs to be paused, abandoned, or resumed.

A.3.7 CSF Data Deduplication Example

1. It is necessary to create a volume (see Figure B.3. 5 Volume Creation) to perform data deduplication.



Figure B.3. 5 Volume Creation

2. Configure data duplication to be enabled on the created volume by sending a configuration parameter to the data deduplication CSF. This operation is illustrated in Figure B.3. 6 Enable Deduplication.




Figure B.3. 6 Enable Deduplication

3. Verify that data deduplication is enabled and ready for the created volume by retrieving status. The status indicates that the configurable parameters have been set. This operation is illustrated by Figure B.3. 7 Retrieve Status.



Figure B.3. 7 Retrieve Status

4. Turn off the schedule of the post process data deduplication on the volume by sending the data deduplication CSF a request that data deduplication not be scheduled to occur on the volume. This is shown in Figure B.3. 8 Disable Deduplication Scheduling.




Figure B.3. 8 Disable Deduplication Scheduling

5. Mount the volume to a server and copy files from existing server directories into the new directory as shown in Figure B.3. 9 Copy files. This writes the files to the newly created

volume and since there is no deduplication scheduled, the data is not deduplicated.

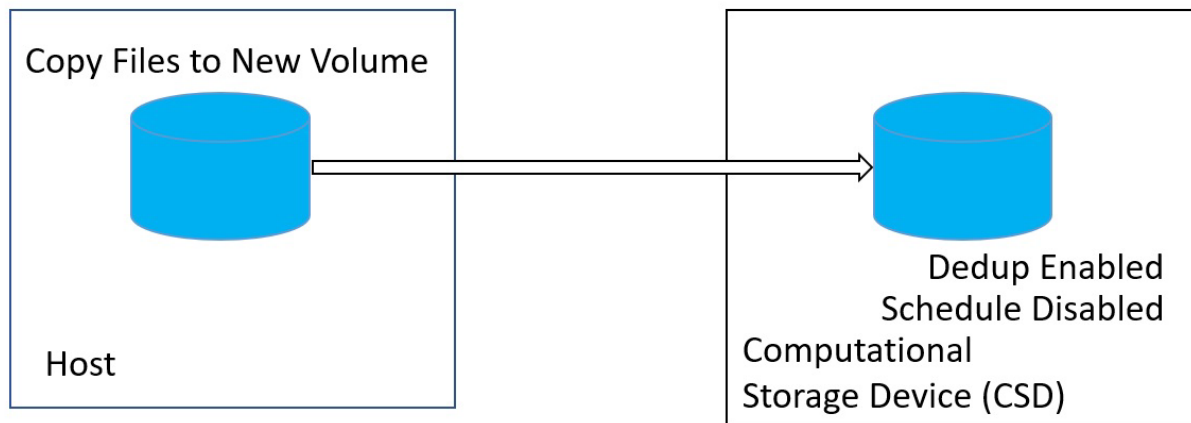


Figure B.3. 9 Copy files

6. Examine the volume for the storage consumed and space using existing server tools as shown in. Figure B.3. 10 Verify space utilized. No space savings have been made since deduplication has yet to occur.

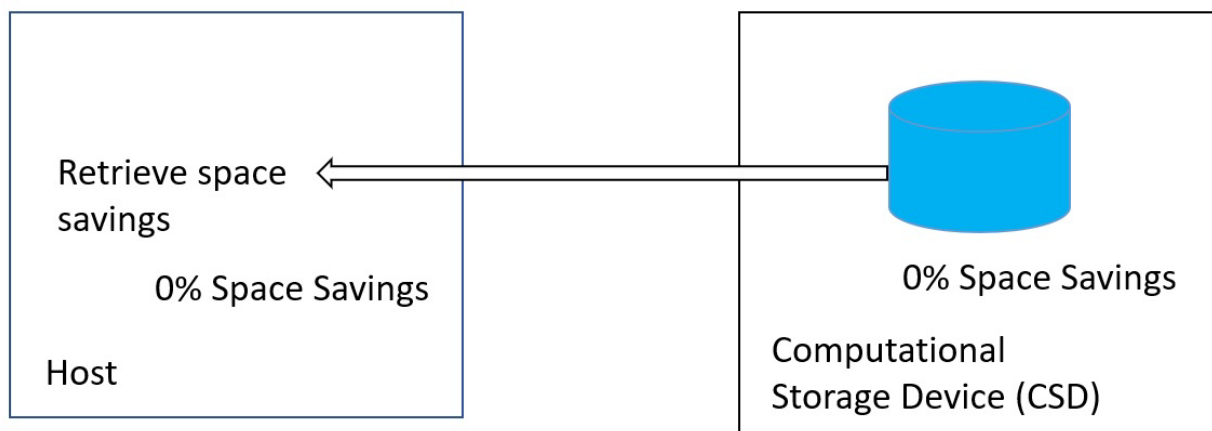


Figure B.3. 10 Verify space utilized

7. Send a request to the data deduplication CSF, to schedule the data deduplication immediately on the volume as shown in Figure B.3. 11 Schedule deduplication. The

data deduplication then occurs on the volume.




Figure B.3. 11 Schedule deduplication

8. Monitor the progress of deduplication by sending the data deduplication CSF a status request. The data deduplication CSF returns the amount of space that has been deduplicated and the percentage complete. This is shown in Figure B.3. 12 Monitor Progress.

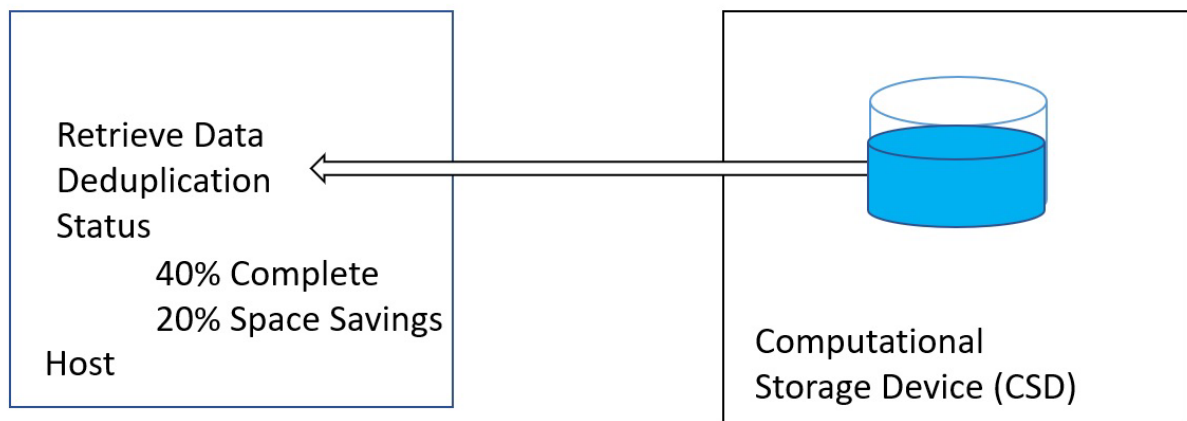


Figure B.3. 12 Monitor Progress

9. When the deduplication is complete, check the space savings by sending a status request to the data deduplication CSF as shown in Figure B.3. 13 verify space savings. The space savings depends on the amount of data that was deduplicated on the

volume.

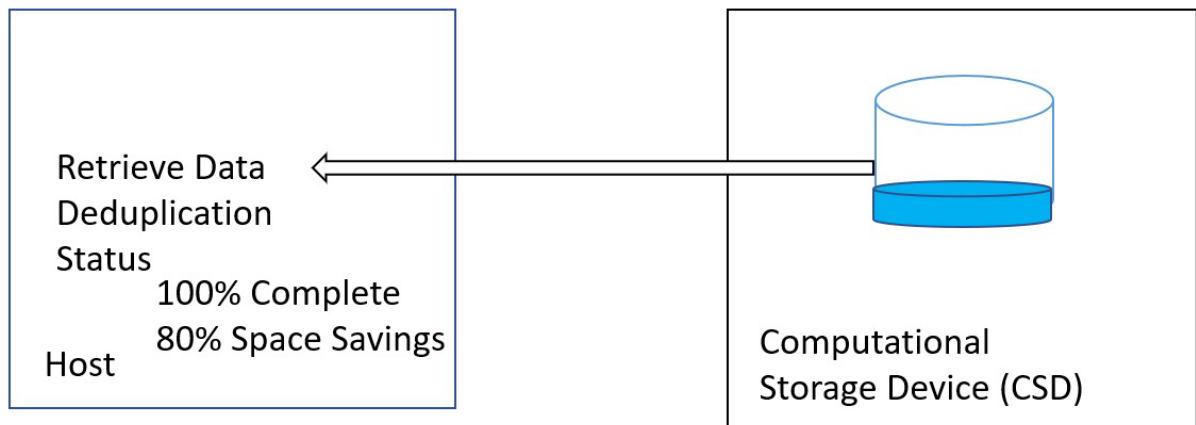


Figure B.3. 13 verify space savings

A.4 A Computational Storage Function on a NVMe Express and OpenCL based Computational Storage Drive Illustrative Example

A.4.1 CSEE Example

A PCIe OpenCL-based Computational Storage Drive (CSD) that consists of an NVMe Controller and an OpenCL accelerator (CSE) which can execute Computational Storage Functions (CSFs). The Computational Storage interface is done over OpenCL not over NVMe. In this example, the NVMe controller and OpenCL accelerator appear as two separate PCI physical functions (PF) on the host. The OpenCL accelerator exposes part of the Function Data Memory (FDM) over the PCIe BAR of its PCI physical function (PF) and this exposed memory is mapped into the host's address space. This enables the host software to allocate buffers from this PCIe BAR memory and use them to move data between the NVMe controller and the OpenCL accelerator directly while bypassing the host system memory entirely. This feature is called PCIe peer-to-peer (P2P) transfer. For the rest of this illustrative example the term P2P is associated with the memory exposed over PCIe BAR and the direct transfer mechanism between the NVMe controller and the OpenCL accelerator. Note that the P2P memory is part of the FDM which is part of the computational storage resources (CSR) and not associated with NVMe storage controller within this CSD

Other specific assumptions for this illustrative example include:

- An application class processor running a modern Linux kernel and user-space distribution.
- A single host system. i.e. no consideration for multi-port NVMe devices.
- No virtualization.
- Leverage existing PCIe and NVMe methods around security.

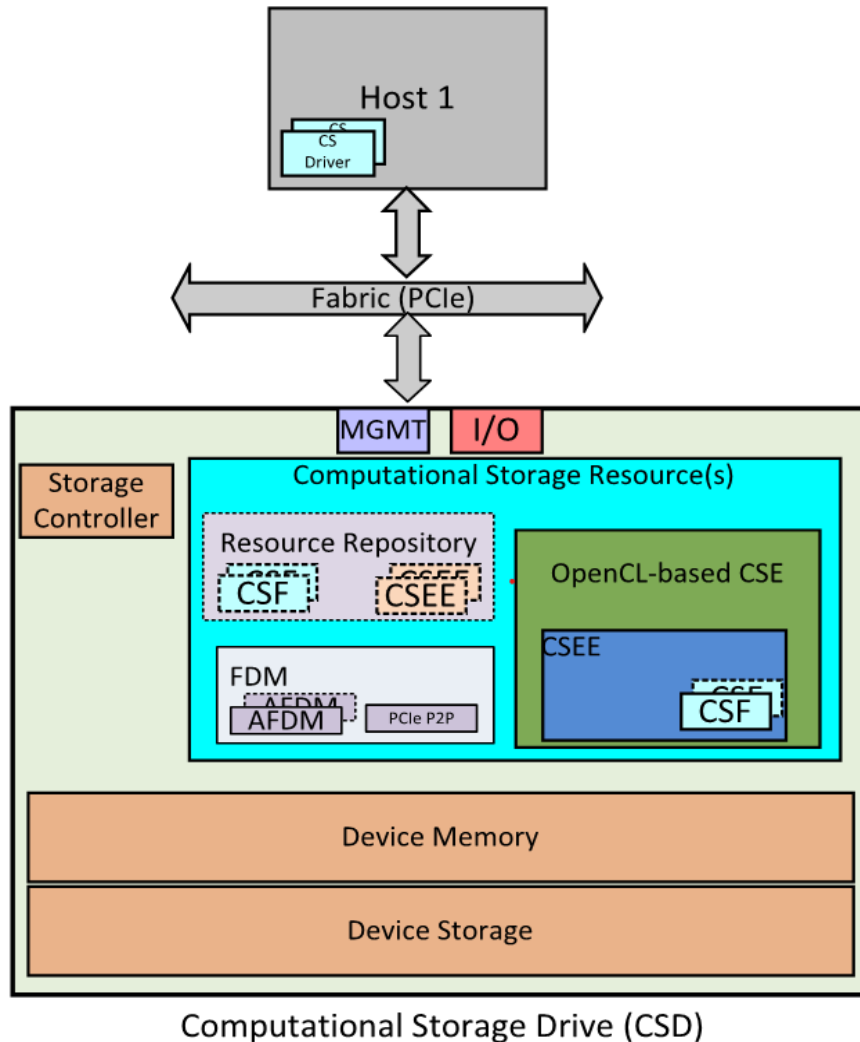


Figure B.4.1 CS architectural diagram adapted for this illustrative example

A.4.2 Computational Storage Drive

Figure B.4.1 describes the elements comprising a CSD. There is a Computational Storage Engine (CSE) combined with persistent memory data storage in the form of an NVMe SSD within a single device (represented by Storage Controller and Device Storage blocks in the figure). The CSEE provides an environment for executing the CSF. If the input data for the CSF is on the device storage, it can be directly transferred (using P2P transfer mechanism) to an AFDM buffer which is allocated from the P2P memory. Similarly, if the output data from the CSF is to be stored on the device storage, it is able to be directly transferred from the P2P AFDM buffer using P2P transfer mechanisms. This example also applies to a case with two discrete PCIe devices connected to the same PCIe fabric (i.e., they are not part of same device as shown in the above

figure but are port of two different physical devices). In this case the OpenCL device is referred to as a Computational Storage Processor(CSP) with an exposed PCIe P2P BAR memory and the second device is an NVMe SSD.

This example does not explicitly address multiple CSDs in one system, but the architecture, including the software drivers for OpenCL and NVMe all support multiple devices in one system.

A.4.3 Theory of Operation

This section walks through the Theory of Operation and applies it in a specific fashion to this illustrative example. The three main phases are:

- Discovery
- Configuration
- Usage

A.4.4 Discovery

In this illustrative example, the storage controller of the CSD is discovered and configured using the procedures described in the NVMe base specification without any changes. The OpenCL accelerator (CSE) is discovered and configured using the OpenCL APIs.

PCIe devices (NVMe controller and OpenCL device) are discovered by the host at power-up via PCIe bus enumeration or at run-time via a hot-plug event and bus rescan. A modern OS like Linux binds the relevant drivers to the discovered PCIe devices. In this case the NVMe driver is bound to the NVMe controller of the CSD and the OpenCL runtime driver is bound to the OpenCL device.

A.4.4.1 NVMe Function Discovery

Once an NVMe controller has been detected, the NVMe driver on the host discovers the capabilities of said controller via the NVMe Identify Admin command. NVMe has the concept of namespaces, which refers to discrete resources behind an NVMe controller and one controller is able to support many namespaces as well as being able to (optionally) create and delete namespaces (via namespace management commands). Therefore:

- 1) PCIe enumeration discovers the NVMe controller.
- 2) The NVMe driver binds to this PCIe device and controller capabilities and namespaces are discovered using standard NVMe controller Identify Admin commands.
- 3) The driver configures the NVMe namespaces using NVMe NS identify commands and makes them available to the storage stack on the host.

In this example the driver will discover one namespace on the NVMe controller of the CSD

- a) Namespace 1: Conventional LBA based storage namespace.

At this point the information about the namespaces on this NVMe- controller of the CSD is able to be displayed with admin tools like nvme-cli. As an example the output of nvme-cli list for such a device might look like:

Node	SN	Model	Namespace Type	Format or SubType	FW Rev
/dev/nvme0n1	nvme1	Vendor A 1	Conventional LBA	512 B + 0 B	1.0

A.4.4.2 OpenCL CSP Function Discovery

As the OpenCL accelerator is on a separate PCI physical function, the first two stages of discovery for the OpenCL CSE mirror those of the NVMe controller:

1. PCIe enumeration is used to discover the OpenCL CSP.
2. The Linux OpenCL Runtime driver binds to this PCIe device

The OpenCL Runtime driver implements platform-specific features that allow applications to query OpenCL devices, device configuration information, and to create OpenCL contexts using one or more devices.

The OpenCL runtime driver also discovers the PCIe P2P BAR memory exposed by the OpenCL CSE and maps it to the host's physical address space.

OpenCL defines set of APIs to discover OpenCL platforms and devices and to configure them. In OpenCL, a platform is defined as the set of OpenCL devices from a vendor that implement OpenCL functionality. The platform vendor also provides the OpenCL runtime drivers to manage the OpenCL devices. A typical usage of OpenCL device involves discovering the platform of the vendor and discovering the devices on that platform.

The OpenCL Runtime driver implements APIs that allow applications to query OpenCL devices and device configuration information.

The APIs relevant to discovering the OpenCL device are described below:

This API returns the number and list of OpenCL platforms found on the system:

```
cl_int clGetPlatformIDs(cl_uint num_entries, cl_platform_id *platforms, cl_uint *num_platforms)
```

This API is used to get the platform specific info of the desired vendor platform:

```
cl_int clGetPlatformInfo(cl_platform_id platform,
    cl_platform_info param_name,
    size_t param_value_size,
    void *param_value,
    size_t *param_value_size_ret)
```

This API is used to find the number of OpenCL devices and their device-ids of the specific vendor platform selected.

```
cl_int clGetDeviceIDs(cl_platform_id platform,
    cl_device_type device_type,
    cl_uint num_entries,
    cl_device_id *devices,
    cl_uint *num_devices)
```

Finally, the following API can be used to iterate over all the discovered devices and find the desired device based on parameter CL_DEVICE_NAME. In this example the OpenCL accelerator inside the CSD will have a unique name that can be matched.

```
cl_int clGetDeviceInfo(cl_device_id device,
    cl_device_info param_name,
    size_t param_value_size,
    void *param_value,
    size_t *param_value_size_ret)
```

A.4.5 Configuration – Explicit Mode

OpenCL provides a rich set of APIs to configure the CSE. These include mechanism to partition a device, create contexts, download OpenCL kernels (CSF), allocate memory (FDM), and more.

A.4.6 Usage – Storage Direct

The OpenCL device can be further optionally partitioned to sub-devices each having its own context and command queue. The following API is used to create sub-devices.

```
cl_int clCreateSubDevices (cl_device_id in_device ,
    const cl_device_partition_property *properties ,
    cl_uint num_devices ,
    cl_device_id *out_devices ,
    cl_uint *num_devices_ret )
```

Once the device is selected, an OpenCL context must be created on that device using the following API

```
cl_context clCreateContext(cl_context_properties *properties,
    cl_uint num_devices,
    const cl_device_id *devices,
```

```

void *pfn_notify (
    const char *errinfo,
    const void *private_info,
    size_t cb,
    void *user_data
),
void *user_data,
cl_int *errcode_ret)

```

The OpenCL context can be equated with the CSEE in the CSD. The context is used to create a command queue between the host and the device and to load OpenCL kernels, allocate memory objects.

The API to create the command queue is:

```

cl_command_queue clCreateCommandQueue(cl_context context,
    cl_device_id device,
    cl_command_queue_properties properties,
    cl_int *errcode_ret)

```

After the command queue is created, an OpenCL program is able to be created from a binary file (whose format is specific to the OpenCL device execution environment). The OpenCL program may contain one or more OpenCL kernels. An OpenCL kernel is equivalent to a CSF on the CSD.

The following APIs are used to build the OpenCL program from binary and create OpenCL kernels (CSFs)

```

cl_program clCreateProgramWithBinary (    cl_context context,
    cl_uint num_devices,
    const cl_device_id *device_list,
    const size_t *lengths,
    const unsigned char **binaries,
    cl_int *binary_status,
    cl_int *errcode_ret)

cl_kernel clCreateKernel ( cl_program program,
    const char *kernel_name,
    cl_int *errcode_ret)

```

The input and output buffers for the OpenCL kernels are allocated on the device and mapped to host side buffers using the following APIs

The following API allocates a buffer on the OpenCL device and a corresponding buffer on the host memory

```

cl_mem clCreateBuffer (cl_context context,

```

```

    cl_mem_flags flags,
    size_t size,
    void *host_ptr,
    cl_int *errcode_ret)

```

The following API returns the host mapped address of the buffer created in the previous API.

```

void * clEnqueueMapBuffer (cl_command_queue command_queue,
    cl_mem buffer,
    cl_bool blocking_map,
    cl_map_flags map_flags,
    size_t offset,
    size_t cb,
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list,
    cl_event *event,
    cl_int *errcode_ret)

```

The device side buffer handle returned by `clCreateBuffer` APIs is passed to OpenCL kernel using the following API

```

cl_int clSetKernelArg (cl_kernel kernel,
    cl_uint arg_index,
    size_t arg_size,
    const void *arg_value)

```

There is no explicit configuration required for the NVMe controller as the host driver already enumerates and configures the namespaces during the discovery process as per the NVMe standard

A.4.7 Usage – Explicit Mode Computational Storage

In this example we look at a simple CSF use case where data stored on the storage of the CSD is read, a transformation is performed on that data and the result is returned to the host.

As the input data is stored on the NVMe data storage of the CSD, it needs to be first read and given to the CSF. There are two ways to do this: The host issues an NVMe read operation and reads the data into its memory. Once that is complete, moves that data to OpenCL device memory (AFDM) using the OpenCL API `clEnqueueMigrateMemObjects`. The other method is to make use of the P2P mechanism. In this mechanism the host issues NVMe read operation but the destination of that operation is the P2P memory on the OpenCL device. Using the PCI P2P operation, the NVMe controller directly moves the data into OpenCL device memory (AFDM). Using the P2P mechanism reduces the extra data movement from NVMe storage to host memory and then back to OpenCL device memory (AFDM).

Once the data is available in AFDM, the OpenCL kernel (CSF) can be triggered to run using the following API.

```
cl_int clEnqueueTask (cl_command_queue command_queue,  
    cl_kernel kernel,  
    cl_uint num_events_in_wait_list,  
    const cl_event *event_wait_list,  
    cl_event *event)
```

The host now waits for the CSF to complete its operations using the API

```
cl_int clWaitForEvents(  
    cl_uint num_events,  
    const cl_event* event_list)
```

Once the operation is complete, the result data can be moved back to host memory using the OpenCL API `clEnqueueMigrateMemObjects` API

```
cl_int clEnqueueMigrateMemObjects ( cl_command_queue command_queue ,  
    cl_uint num_mem_objects ,  
    const cl_mem *mem_objects ,  
    cl_mem_migration_flags flags ,  
    cl_uint num_events_in_wait_list ,  
    const cl_event *event_wait_list ,  
    cl_event *event )
```

If the use case is to store the result of the CSF back to the NVMe storage of the CSD, the output buffer also can be allocated from the P2P area in the AFDM and NVMe write operation can be initiated to move it into the storage directly using the P2P transfer mechanism between NVMe controller and the OpenCL device.

A.5 Data Compression CSF Example

A.5.1 Overview

Data compression aims to reduce the amount of data stored on a storage device. By reducing the amount of data being physically written to or read from storage media (e.g., flash memory chips), data compression can also improve the storage device IO performance and lifetime span, in addition to storage cost saving.

Different from off-loading compression/decompression through a specific API, Data Compression CSF carries out compression/decompression on the IO data path, transparently from the host (i.e., host simply issues normal IO write/read requests to a CSD with a Data Compression CSF, without calling any other specific API). In order to materialize the storage cost reduction, Data Compression CSF must expose a logical storage space that is larger than its internal physical storage space. However, due to the runtime variation of data compressibility, it is possible that the physical storage space will be used up before the logical storage space is used up. Moreover, it is desirable for users to know how well different files or objects are compressed by Data Compression CSF. Therefore, Data Compression FCSS must provide adequate reporting mechanism and observability in terms of data compression.

A.5.2 Theory of Operation

In order to successfully implement data compression CSF that operates on all write data, a CSx must be able to communicate that it contains a CSE that is able to perform data compression CSF. Once the CSF's ability is determined, the CSF should allow the host to monitor and query the effect of data compression. This illustrative example will attempt to provide the initial framework of the following steps needed to successfully allow for a CSx to perform a data compression CSF.

A.5.3 Discovery

In order to determine if a CSx supports the Data Compression CSF, the CSx must first be discovered as a CSx, with a CSE that is able to perform the function. If the CSE is able to perform a Data Compression CSF, then the CSE must also indicate if the Data Compression CSF allows configuration. If the CSE allows for configuration parameters to the Data Compression CSF, then the configurable parameters must be shared. The following is a possible list of configurable parameters:

1. Supported compression block sizes – The acceptable values of compression data block size. The block size can be set globally for the entire Data Compression CSF, be set separately for each logical storage space region of the Data Compression CSF, or even be set for each individual write request sent to the Data Compression CSF.
2. Supported maximum logical storage space – The maximum size of the logical storage space that can be exposed by the Data Compression CSF.

3. Monitoring – The type of information to collect during runtime operation. Such information can include the runtime physical storage capacity usage of the entire Data Compression CSF, and the runtime physical storage capacity usage of any given logical storage space.

Note that discovery may also return the default values of the configurable parameters.

A.5.4 Configuration

Once the configurable parameters of a CSE that supports the Data Compression CSF are discovered, the Data Compression CSF is able to be configured if allowed. It is possible that the Data Compression CSF will have default values that the user may not want to override, and configuration is not necessary. If the user wants to configure the Data Compression CSF, then the parameters returned by discovery need to be set and sent to the Data Compression CSF before the data compression operation is performed.

A.5.5 Monitoring

As the data is being written to the Data Compression CSF, the user can get statistics on the compression. The following are possible statistics to monitor:

1. Current physical storage space usage of the entire CSF
2. Current physical storage space usage of any logical storage space region of CSF
3. Lifetime data compression ratio of all the data written to CSF so far
4. Data compression ratio over a specified amount of data that has been written to the CSF

A.6 Data Filter CSF Example

A.6.1 Overview

As one important category of operations in data analysis, data filtering aims to filter out the data that are not needed by a query. In conventional practice, a CPU (or GPU) is responsible for data filtering, which requires transferring all the raw data from the storage device into the CPU (or GPU) memory. The unique feature of Data Filter CSF is to pushdown the data filtering operation from the CPU (or GPU) to a CSE in the CSD. It can offload the data filtering operation from the CPU (or GPU), leading to higher system performance, and less host resource contention in terms of CPU/GPU cycles, memory capacity and bandwidth, and I/O bus bandwidth. The following example illustrates using data filter CSF to carry out in-storage data filtering, as shown in Figure 1. A table with four columns is stored in the Data Filter CSF that receives a request “SELECT ID where State=CA” that seeks the IDs of all the table entries in which State equals to CA. As illustrated in the figure, the Data Filter CSF fetch all the table entries from the storage media, extracts each table entry, and checks whether the 4th column in the entry equals to CA. After scanning the entire table, the Data Filter CSF returns the IDs of all the matching entries.

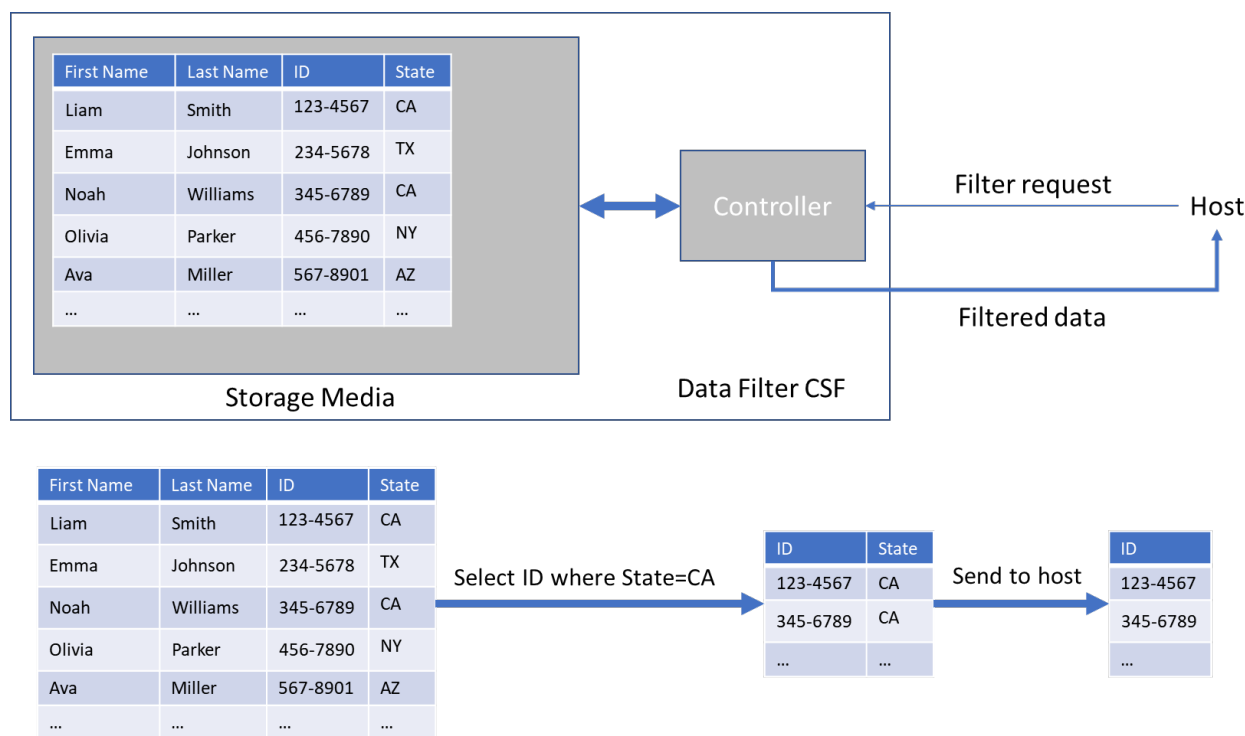


Figure B.6. 1 An example to illustrate the function of data filter FCSS

A.6.2 Theory of Operation

In order to successfully implement data filter as a CSF, a CSx must be able to communicate that it contains a CSE with the ability to perform data filter CSF. Once the ability is determined, the CSF should allow the host to query the supported filter functions and data schema. This illustrative example provides the initial framework of the following steps needed to successfully allow for a CSE to perform data filter as an CSF.

A.6.3 Discovery

The steps in discovery are:

- a) Discover if the Data Filter CSF allows configuration.
 - a. If the Data Filter CSF allows configuration, then discover the configurable parameters. Discovery may also return the default values of the configurable parameters.
 - b. The following is list of possible parameters:
 - i. Supported data formats (e.g. Parquet/ORC, JSON, XML formats). In order to perform data filtering, Data Filter CSF must be able to understand the data format specified by the host.
 - ii. Supported data types (e.g., ASCII string, integer)
 - iii. Supported filtering operations (e.g. >, <, =)
 - iv. Failover – The action to take (e.g. rollback to host-based data filter) if data filter is interrupted.
- b) Discover operational attributes of the Data Filter CSF
 - a. Monitoring Capabilities – The type of process information that can be collected during data filtering. Such possible information includes the I/O rate of the data filtering operation, the size of the data processed, the size of the data remaining to be processed, the percentage complete, etc.
 - b. Operational Interruption Capability – Whether or not the data filtering operation can be interrupted for purpose of either suspending, abandoning, or resuming the request.

A.6.4 Configuration

Once a CSE that supports the Data Filter CSF is discovered, the Data Filter CSF is able to be configured if allowed. It is possible that the Data Filter CSF will have default values that the user may not want to override, and configuration is not necessary. If the user wants to configure the Data Filter CSF, then the parameters returned by discovery need to be set and sent to the Data Filter CSF before the data filtering operation is performed.

A.6.5 Operation

To utilize data filter CSF to carry out in-storage data filtering, host passes enough information about the data filter operation to the CSE that executes the data filter CSF. Accordingly, the data filter CSF performs in-storage data filtering and returns the results back to the host. The information about the data filter operation may include:

1. The address of to-be-processed data
2. The data format (e.g., MySQL, Parquet) and schema (e.g., the number of columns in the table, and data type of each column)
3. The specific filter operation to be performed
4. The host memory address for the returned data

A.6.6 Monitoring

As the data object is being processed by the data filter CSF, the user can get status and statistics on the process. The following are possible status and statistics to monitor (not all items are *required*; not all items are unique to this CSF):

8. The I/O rate of the data filter operation
9. Current amount of data processed by the data filtering operation
10. Current amount of data remaining to be processed by the data filtering operation
11. The percentage of completion of the data filtering operation
12. Success or failure of the operation
13. Existing state of the operation such as paused, interrupted, resumed, etc.

Based on the status and statistics, the user can then determine if the operation needs to be paused, abandoned or resumed.

